


pyGAM

balancing predictive power
and interpretability
with generalized additive models

dani servén marín

Who are you?

- twin
- Spain + USA
- <3
 - Full-stack data (ml, r+d, deploying...)
 - Society + environment
 - learning
 - open source
 - d.i.y
- core contributor to pyGAM

github.com/dswah/

What's pyGAM?

New open-source library for estimating Generalized Additive Models.

- Modular
- Pure Python
- Plays nice with PyData tools

github.com/dswah/pygam

Compatible with sklearn

```
# do regression
gam = LinearGAM().fit(X, y)
y_pred = gam.predict(X)

# or do classification
lgam = LogisticGAM().fit(X, y)
lgam.accuracy(X, y)

# build a custom model
custom_gam = GAM(link='log', distribution='poisson').fit(X, y)
```

Why did you make pyGAM?

- No P-spline GAMs in Python
- Contribute to PyData stack

Why did you make pyGAM?

Human learning:

- Statistics
- Code quality
- Open-source

Not enough time...

Regression

Histogram smoothing

Confidence intervals

Weighted data

Prediction intervals

Sklearn RandomizedSearchCV

Partial dependence

Alternate optimization objectives

Model statistics

Optimizer Callbacks

Classification

Expectiles

Custom GAMs

Sampling from fitted models

Extrapolation

Monotonic / Convex / Concave constraints

Not enough time...

Regression

Histogram smoothing

Confidence intervals

Weighted data

Prediction intervals

Sklearn RandomizedSearchCV

Partial dependence

Alternate optimization objectives

Model statistics

Optimizer Callbacks

Classification

Expectiles

Custom GAMs

Sampling from fitted models

Extrapolation

Monotonic / Convex / Concave constraints

Welcome to PyDatapolis

Need to generate power for our self-sufficient city

- Plan tomorrow's hourly power demand
- too low => blackouts
- too high => waste resources
- Data Driven™

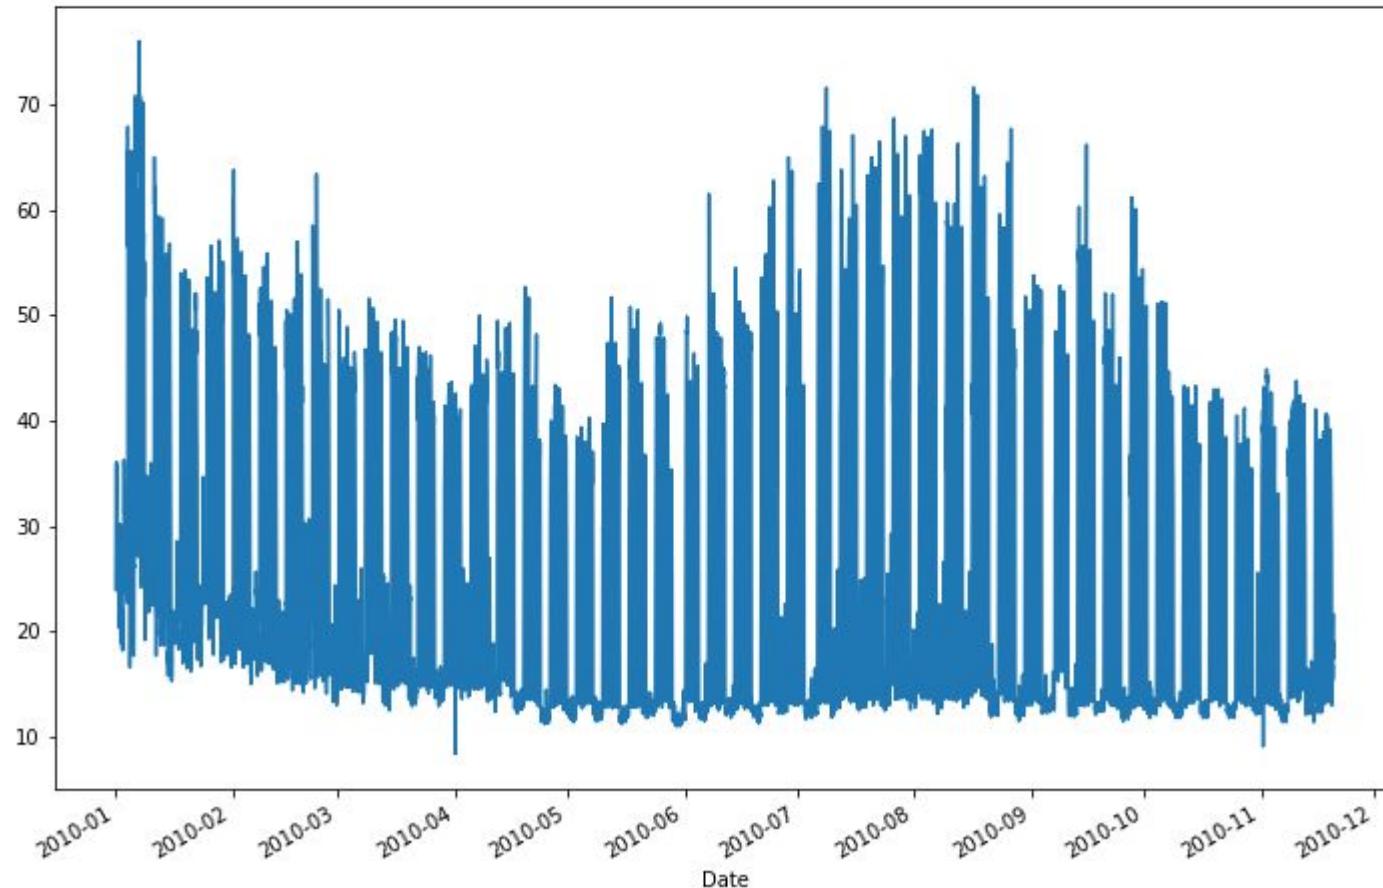
Welcome to PyDatapolis

```
df.columns
```

```
Index(['OAT', 'Building 6 kW', 'Cloud Cover'], dtype='object')
```

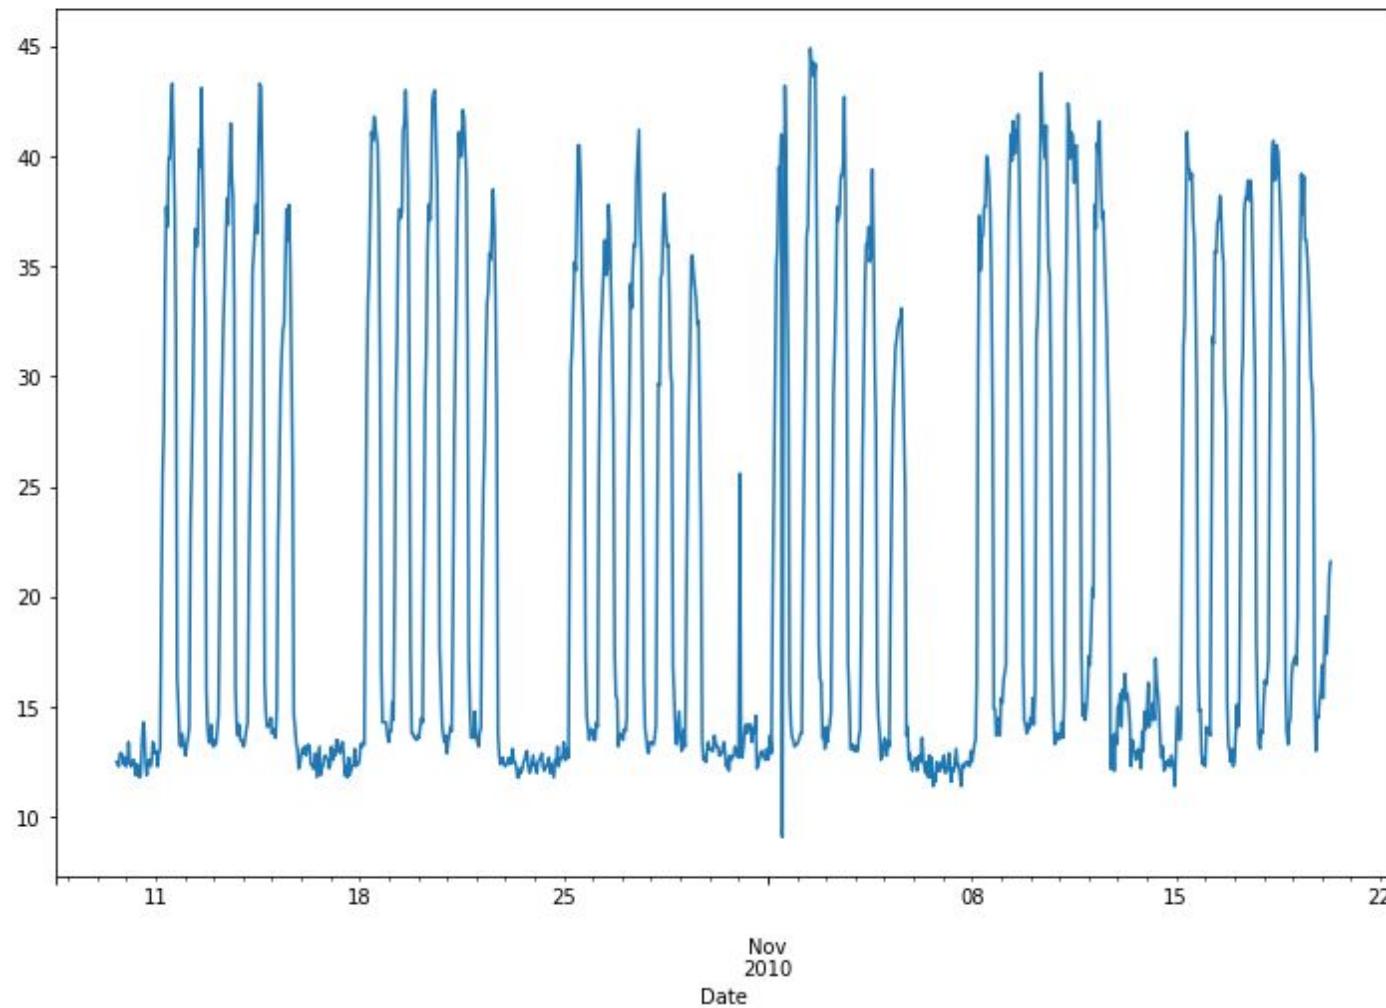
```
df['Building 6 kW'].plot(figsize=(12,8))
```

```
<matplotlib.axes. subplots.AxesSubplot at 0x7f8aa72d57b8>
```



Welcome to PyDatapolis

```
df['Building 6 kw'][-1000:].plot(figsize=(12,8))  
<matplotlib.axes._subplots.AxesSubplot at 0x7feeaac56b00>
```



Welcome to PyDatapolis

- Build LSTM / ARIMAX / kNN / ... model
- Good fit
- Show your comrades the prediction for tomorrow
- But they are skeptical...

Welcome to PyDatapolis

- How certain are you?
- It's going to be a hot summer
 - how should we prepare?...
- What happens on sunny days?
- etc?

Welcome to PyDatapolis

You survey your tool box...

- Linear models have too much bias
- Powerful models are black boxes
- ARIMAX just gives you coefficients
- ...

Welcome to PyDatapolis

You take a break...

read a blog post about GAMs

Welcome to PyDatapolis

You take a break...

read a blog post about GAMs

so you try it out...

```
pip install pygam
```

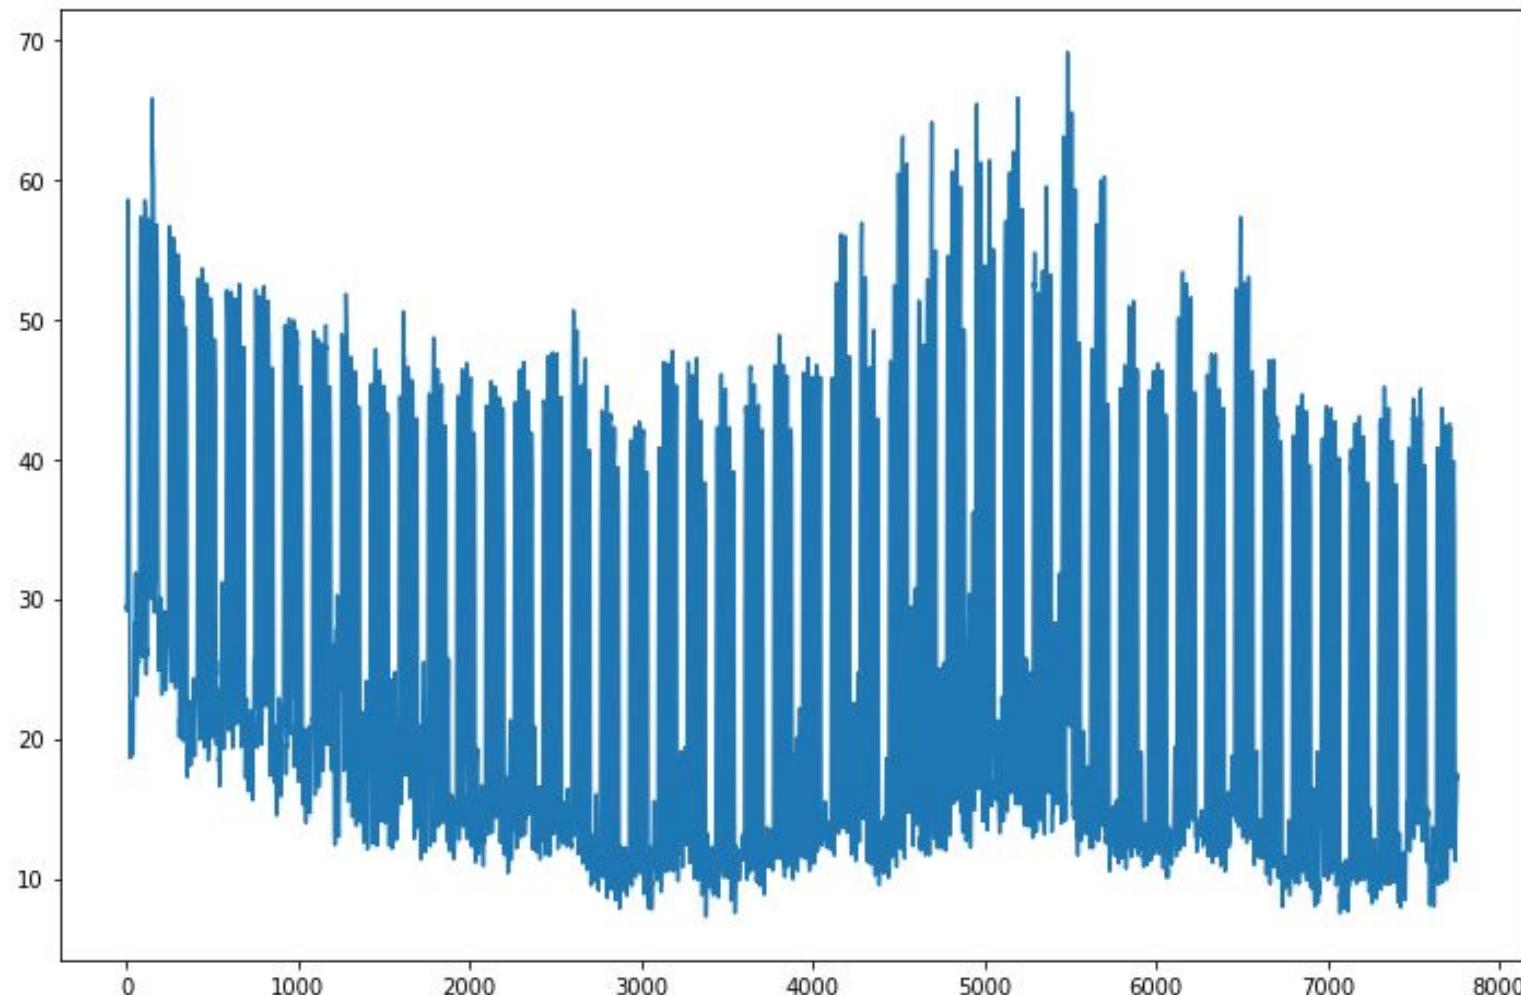
Welcome to PyDatapolis

```
X = df[['day_of_week', 'time_of_year', 'OAT', 'Cloud Cover']].values  
y = df['Building 6 kW'].values  
  
from pygam import LinearGAM  
gam = LinearGAM().fit(X, y)
```

Welcome to PyDatapolis

```
plt.plot(gam.predict(X))
```

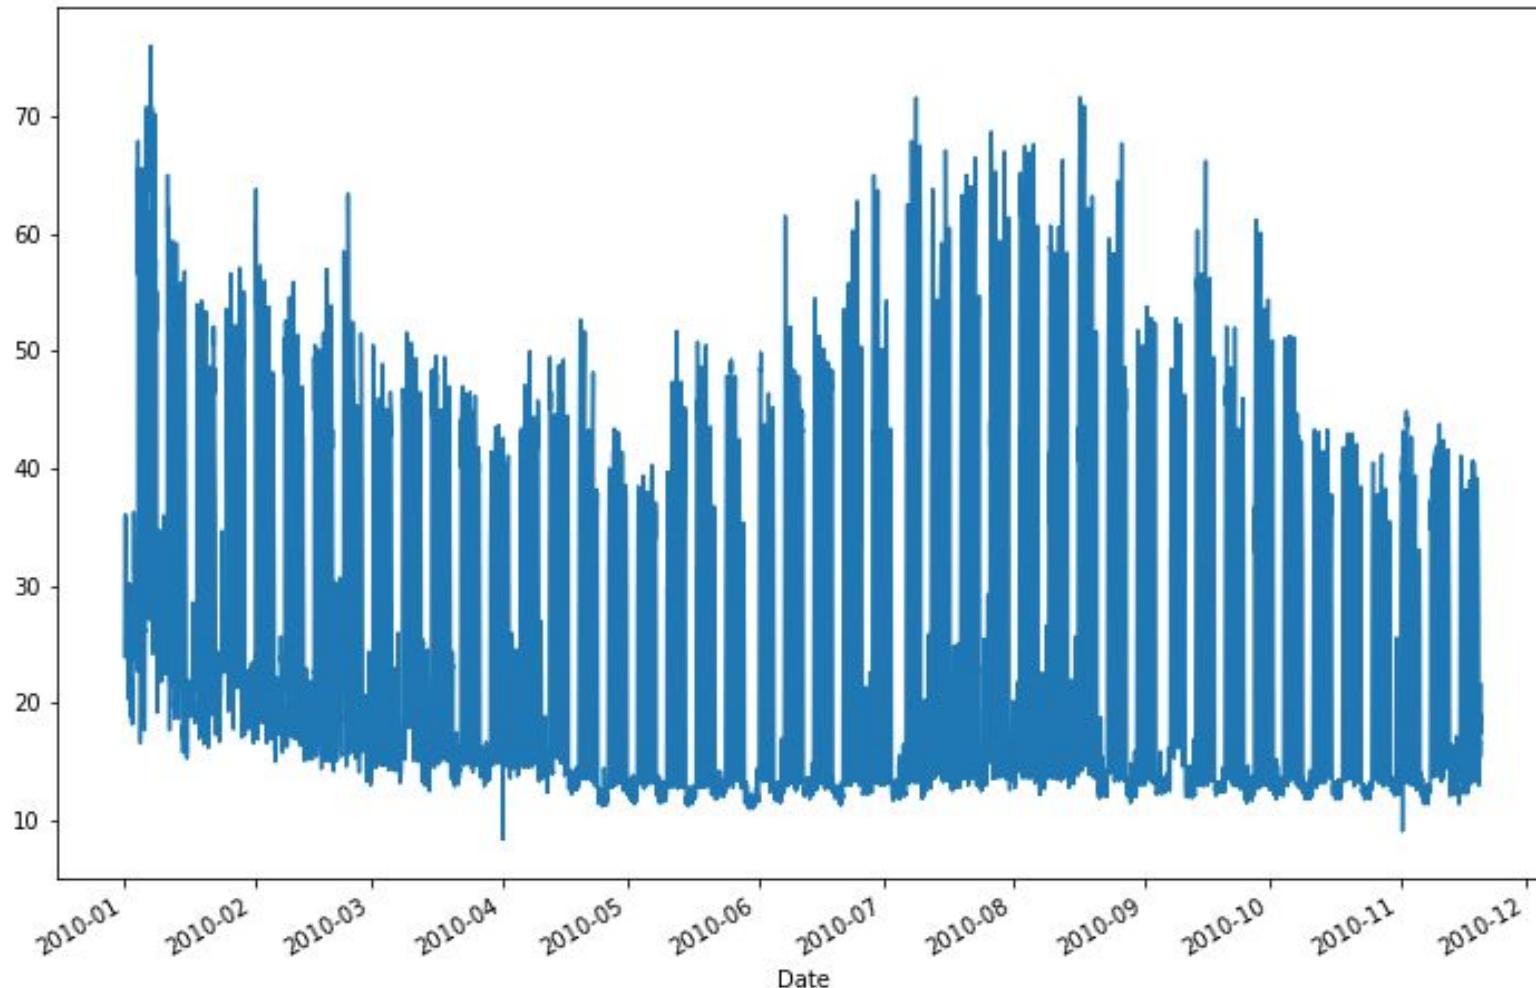
```
[<matplotlib.lines.Line2D at 0x7feeaa8128ba8>]
```



Welcome to PyDatapolis

```
df['Building 6 kw'].plot(figsize=(12,8))
```

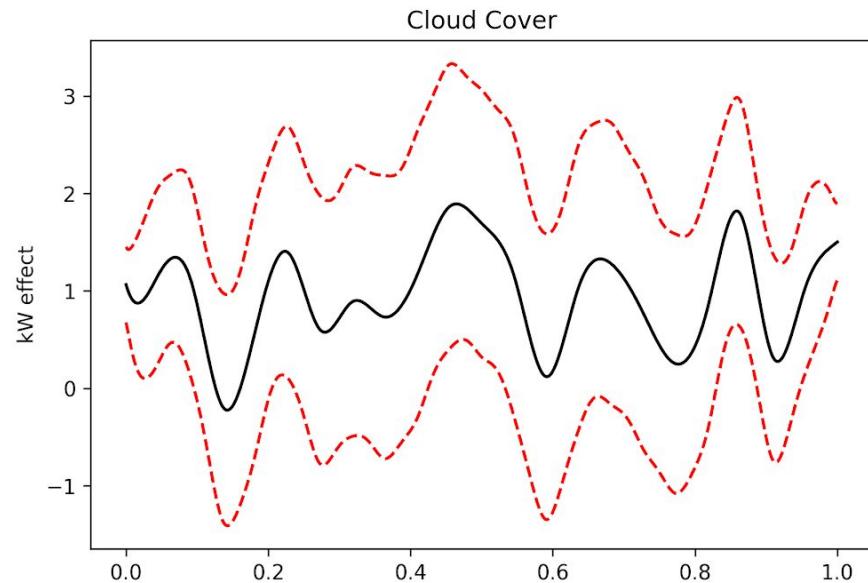
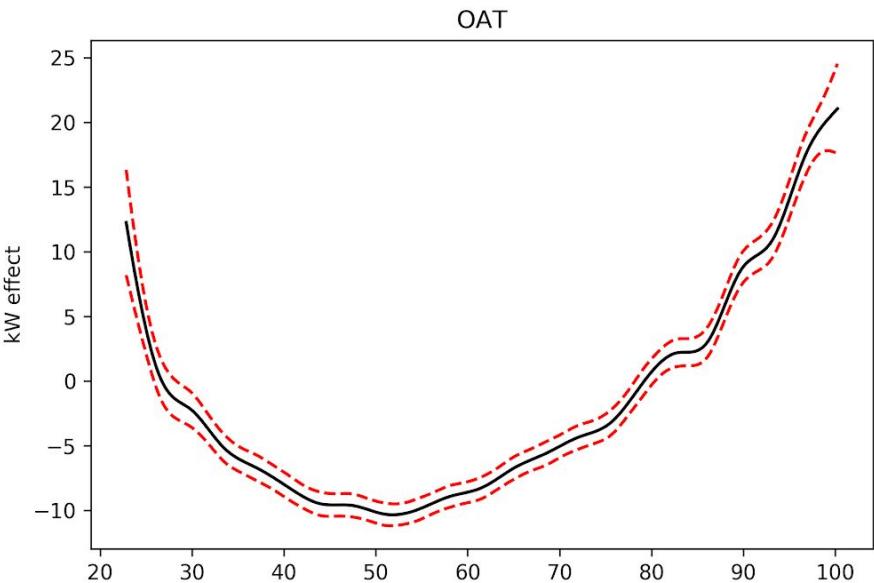
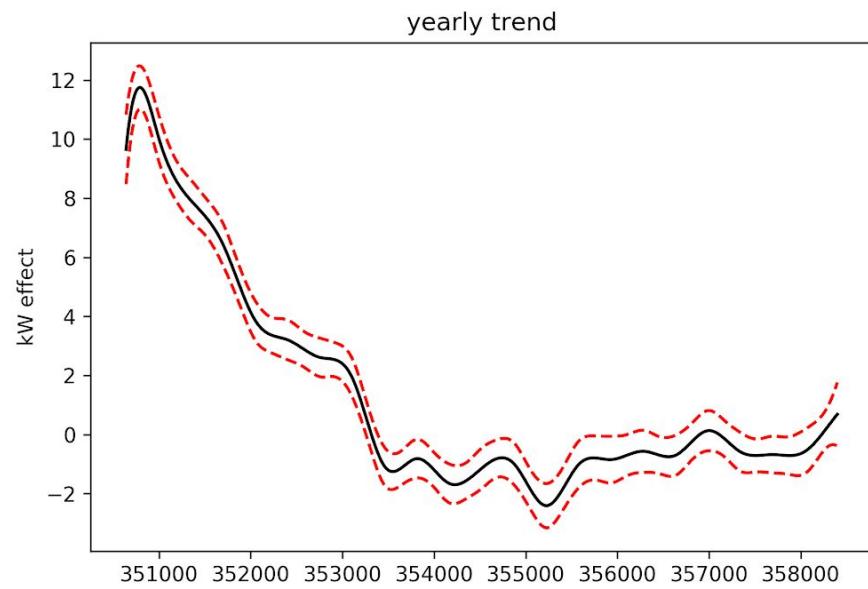
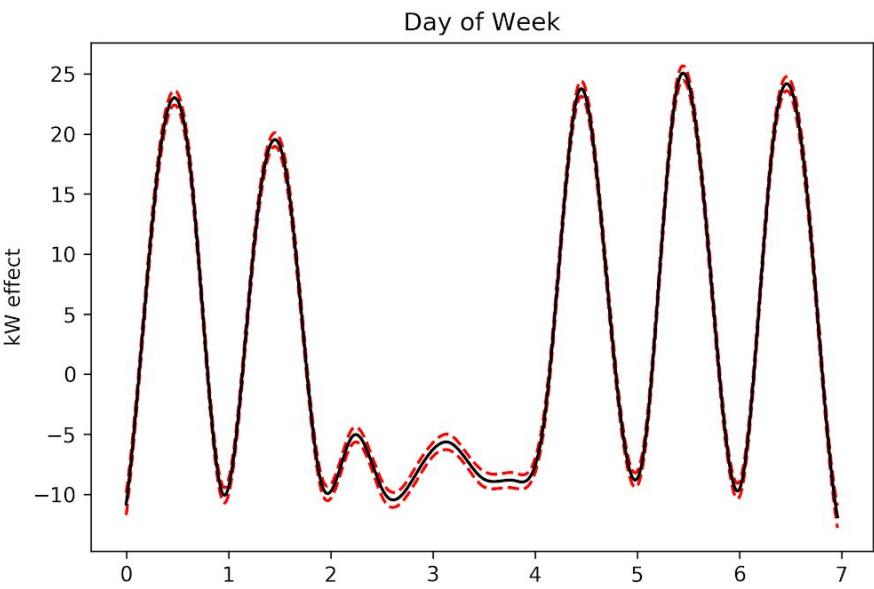
```
<matplotlib.axes. subplots.AxesSubplot at 0x7f8aa72d57b8>
```



What's pyGAM?

```
pdeps, confi = gam.partial_dependence(width=0.95)
```

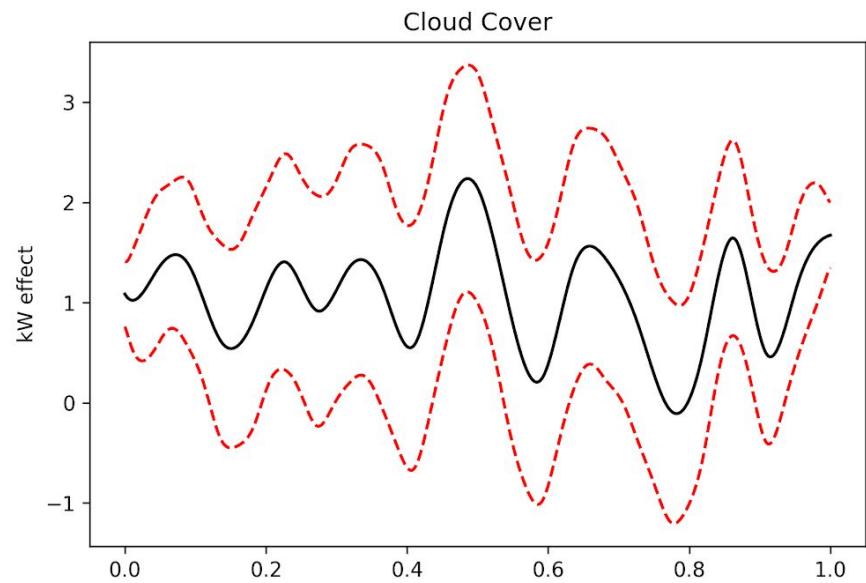
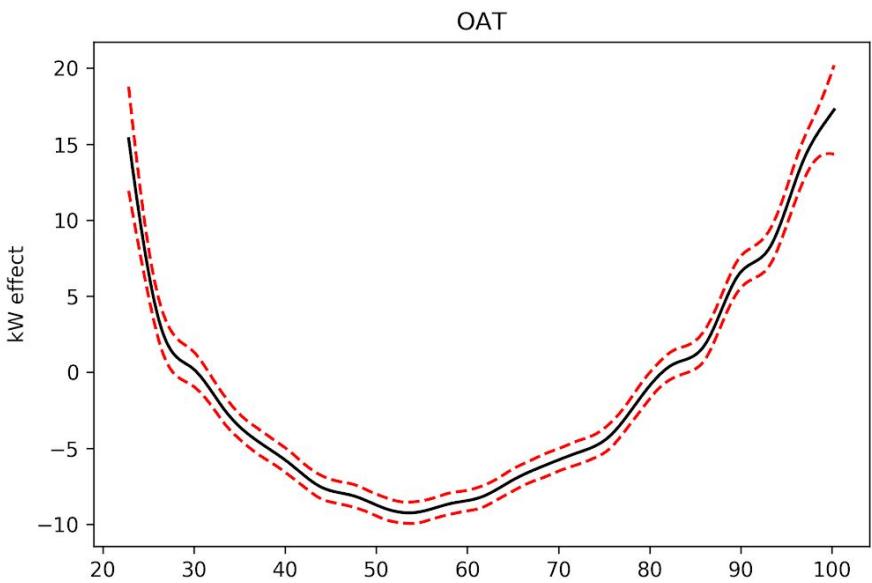
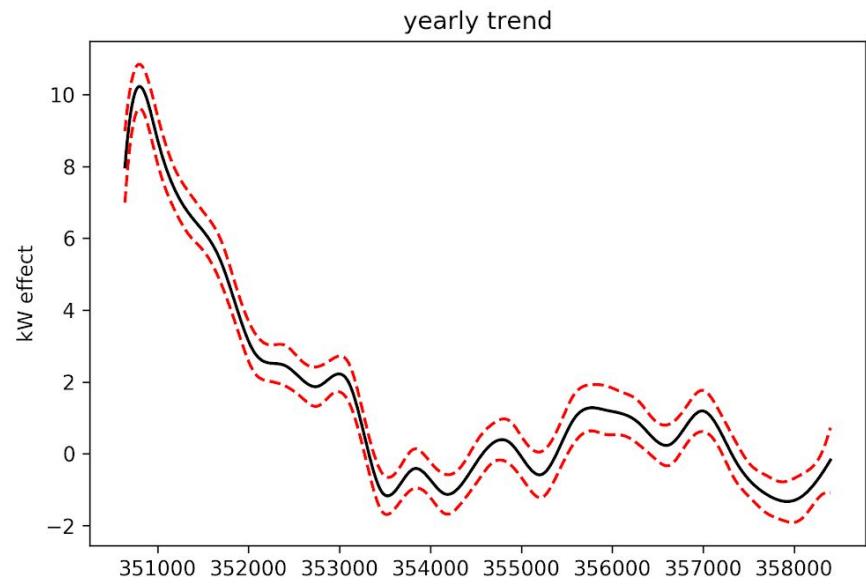
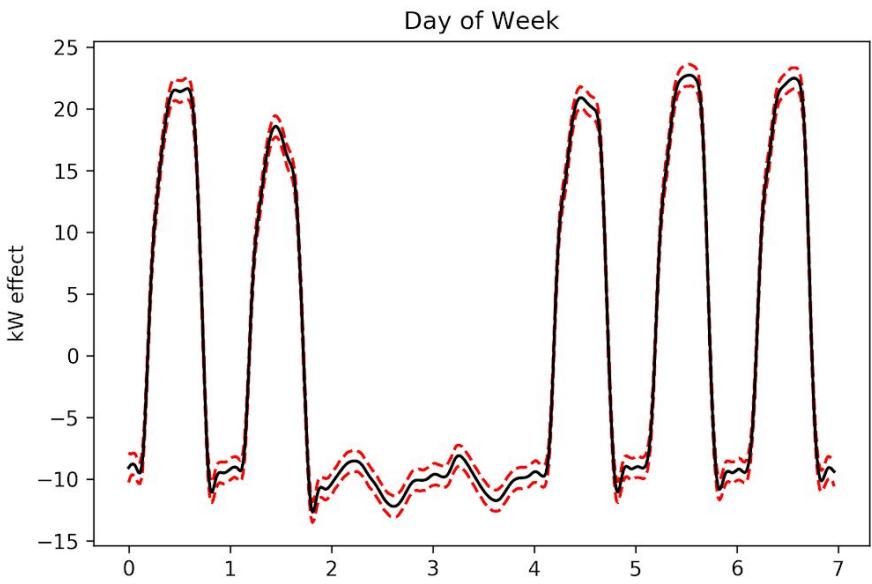
Welcome to PyDatapolis



Welcome to PyDatapolis

```
gam = LinearGAM(n_splines=[100, 25, 25, 25]).fit(X, y)
```

Welcome to PyDatapolis

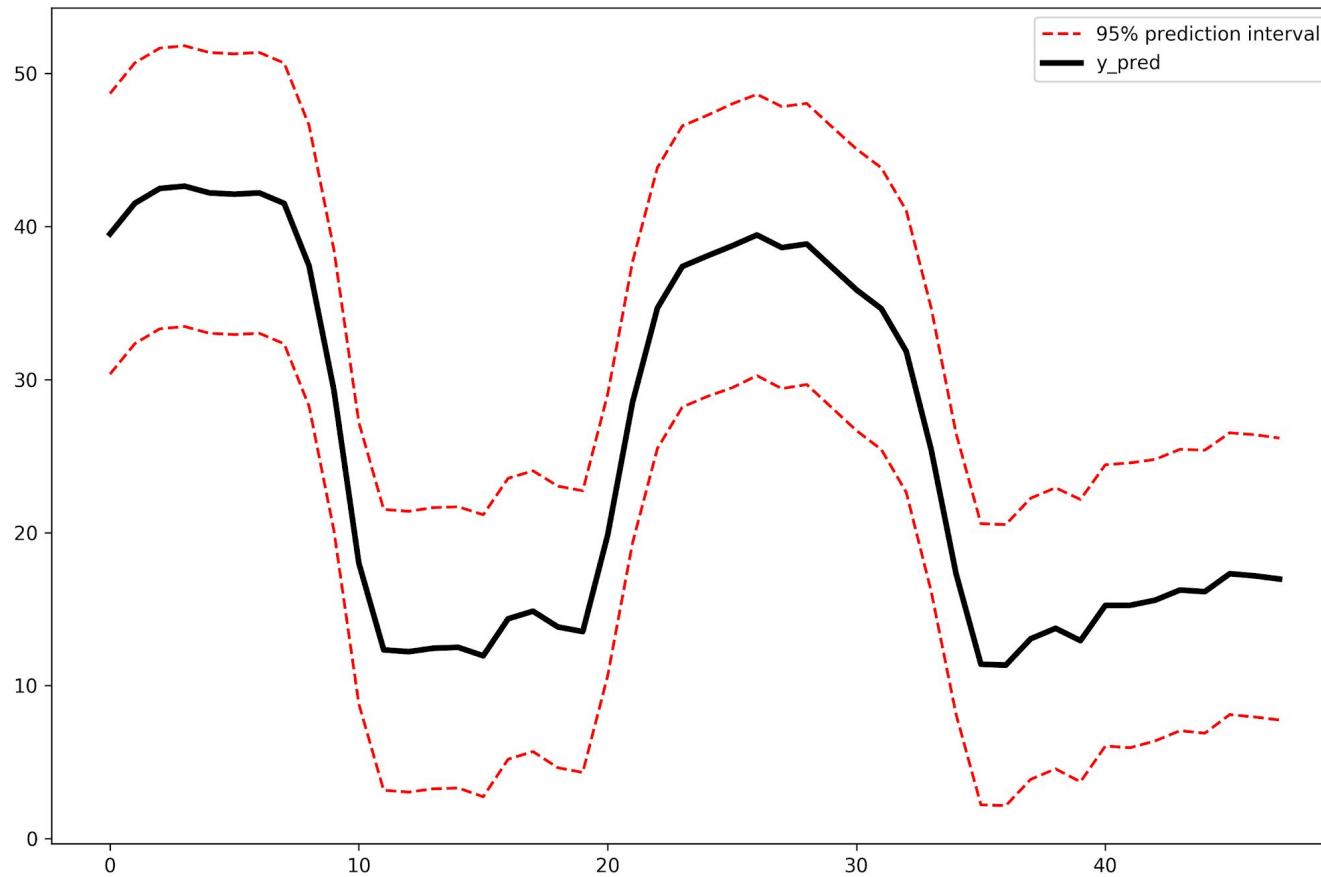


Welcome to PyDatapolis

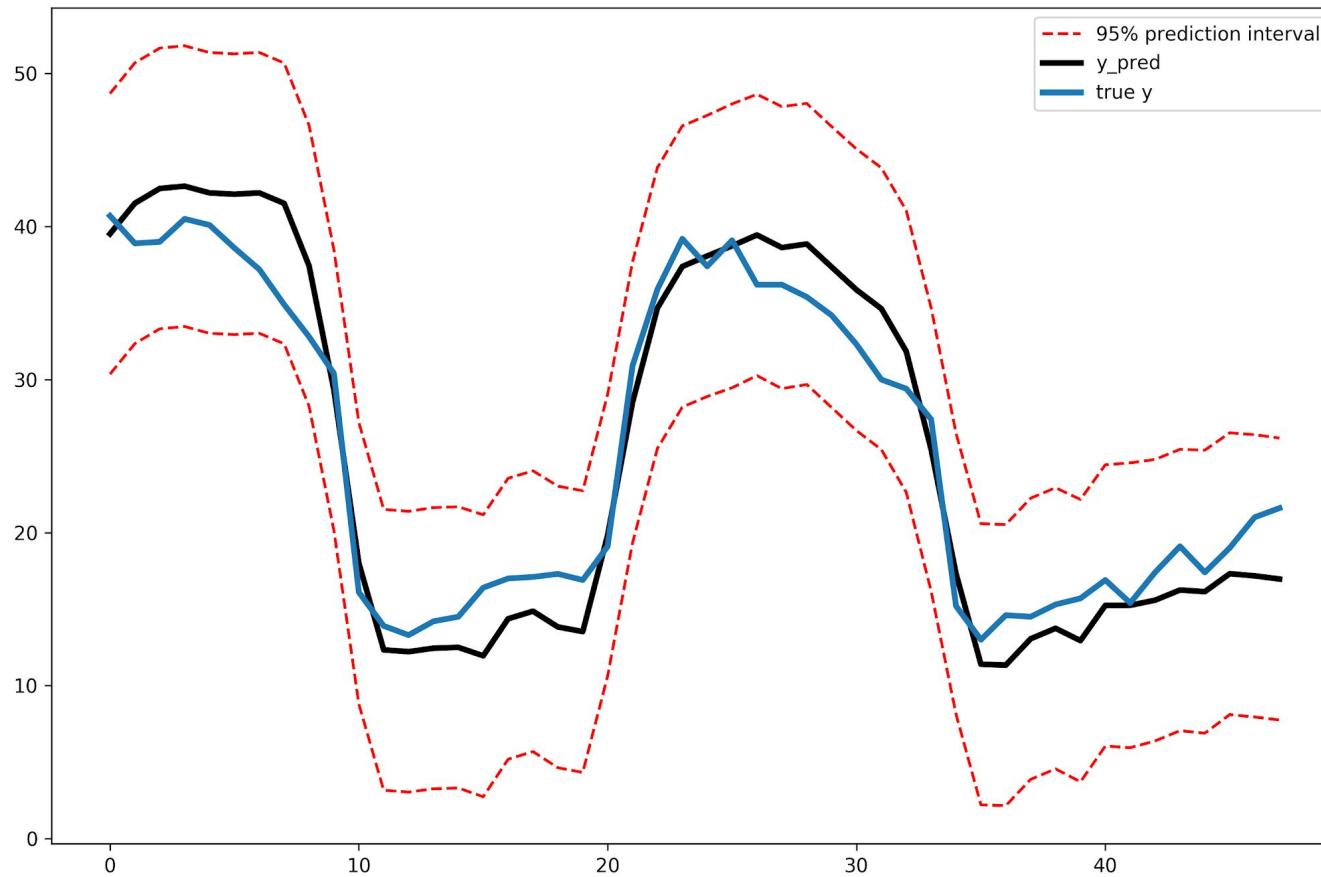
```
y_pred = gam.predict(X_pred)
predi = gam.prediction_intervals(X_pred, width=0.95)

plt.plot(predi, c='r', ls='--')
plt.plot(y_pred, c='k')
```

Welcome to PyDatapolis



Welcome to PyDatapolis



Welcome to PyDatapolis

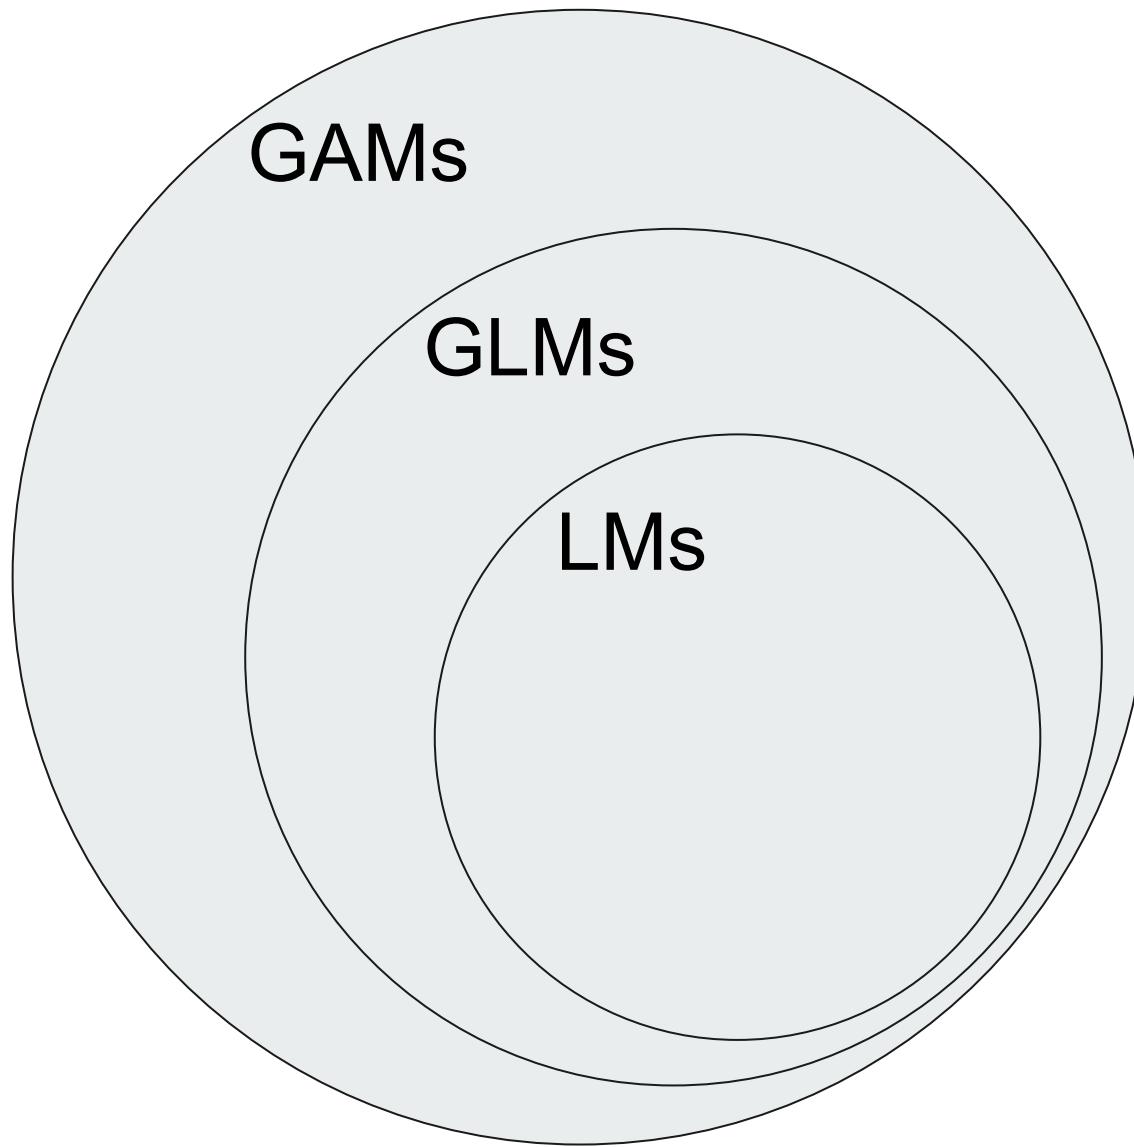
Open questions...

- How do we choose the smoothness?
- Feature selection?

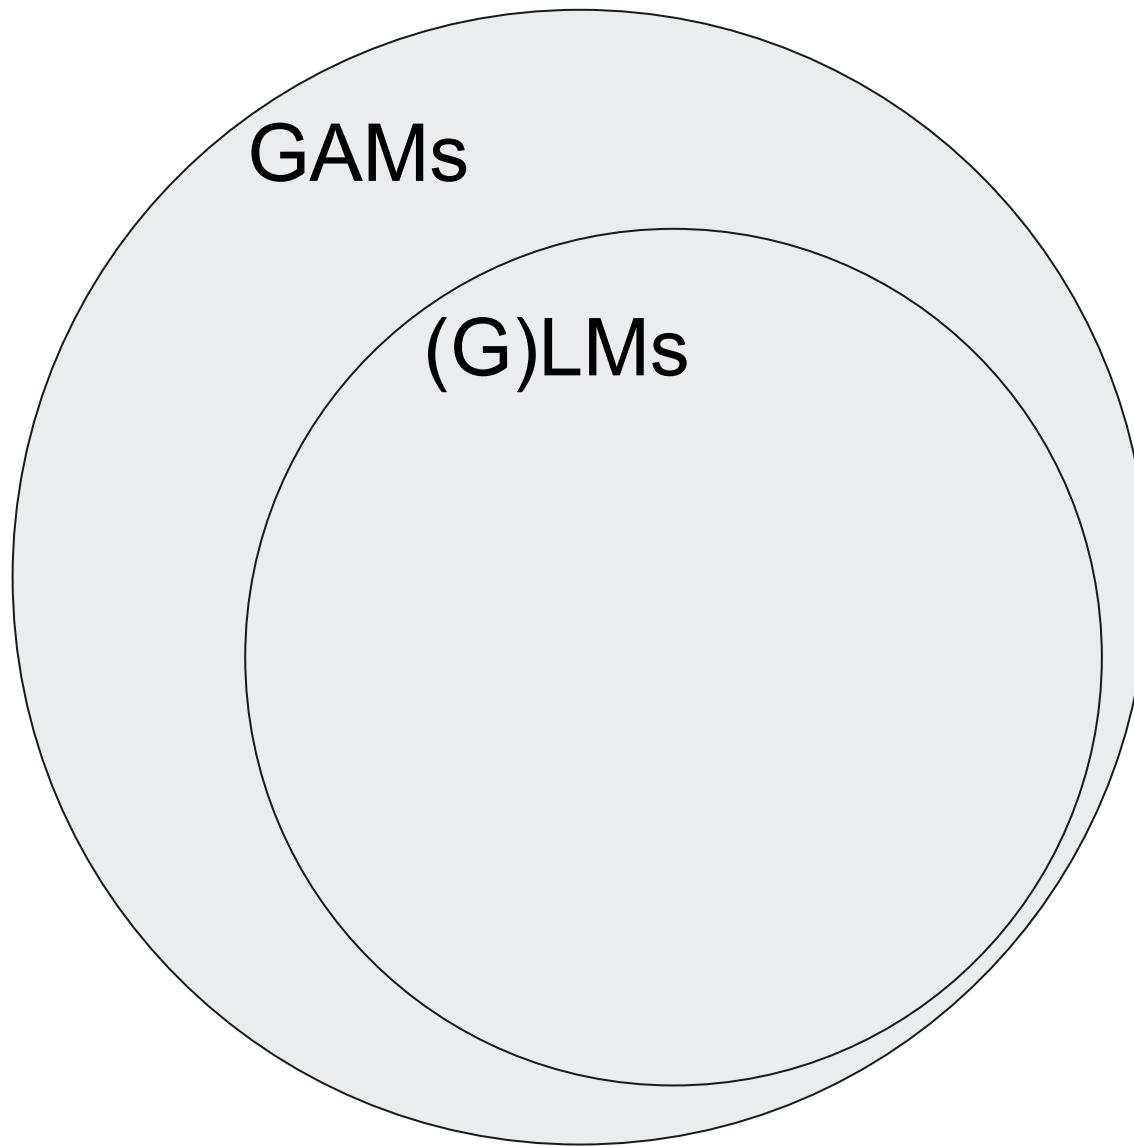
Plan

- What are GAMs
- A brief tour or splines
- Smoothing

What are GAMs

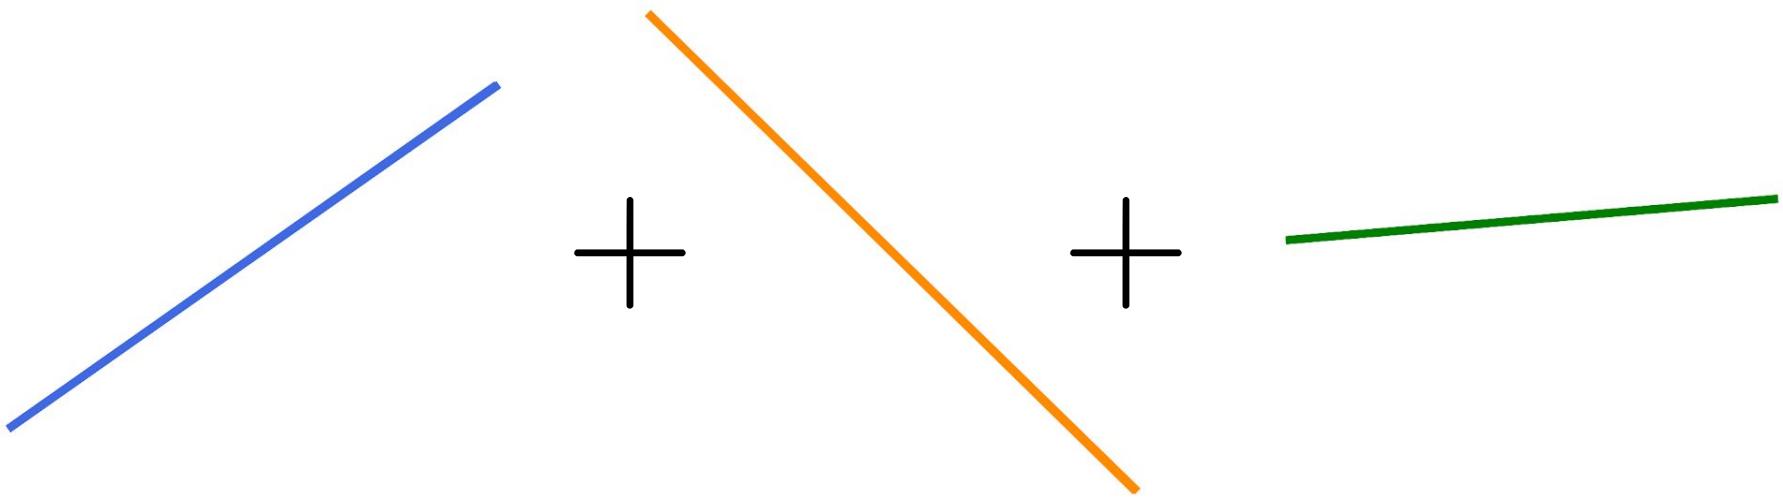


What are GAMs



What are GAMs

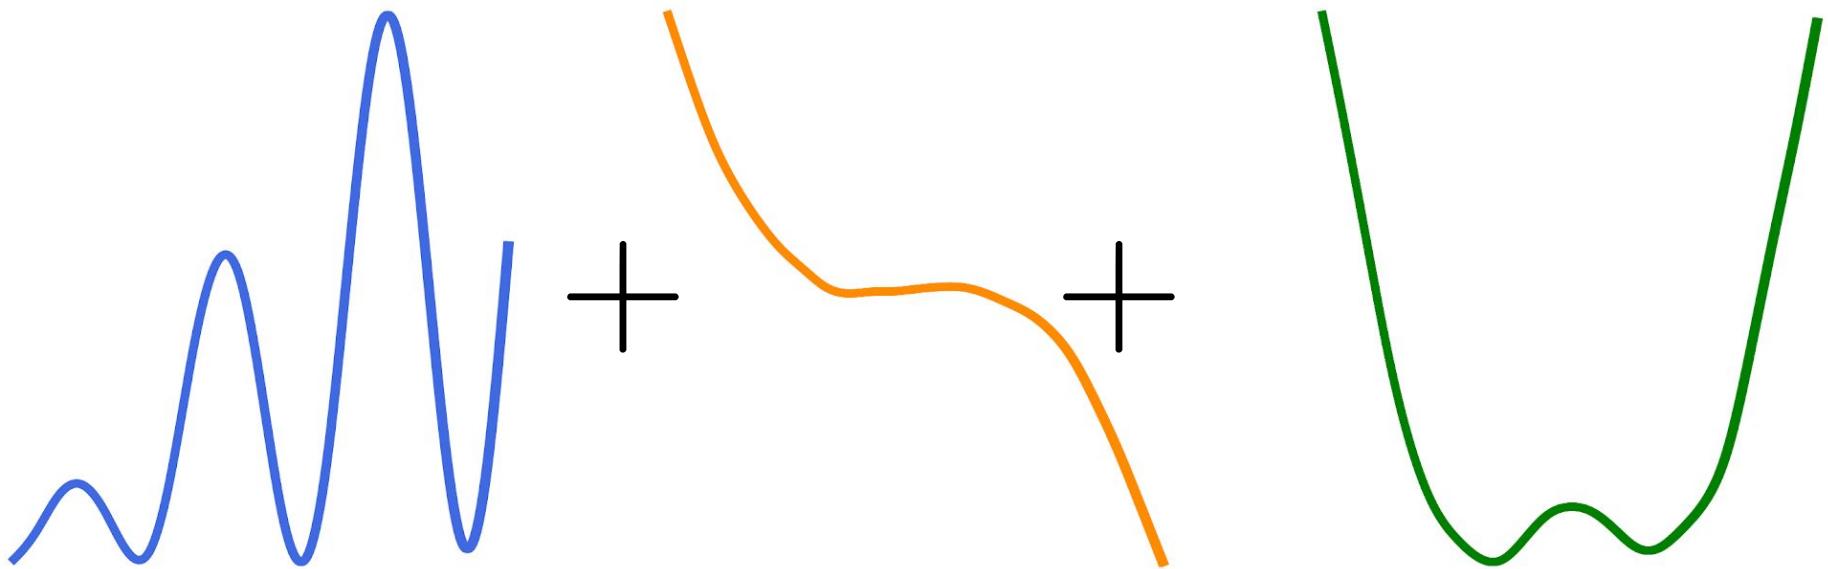
Linear Model



$$y = \beta_1 x_1 + \beta_2 x_2 + \beta_3 x_3 + c$$

What are GAMs

GAM



$$y = f_1(x_1) + f_2(x_2) + f_3(x_3) + c$$

What are GAMs

In general...

$$y = f_1(x_1) + f_2(x_2, x_3) + f_3(x_4) + \cdots + c$$

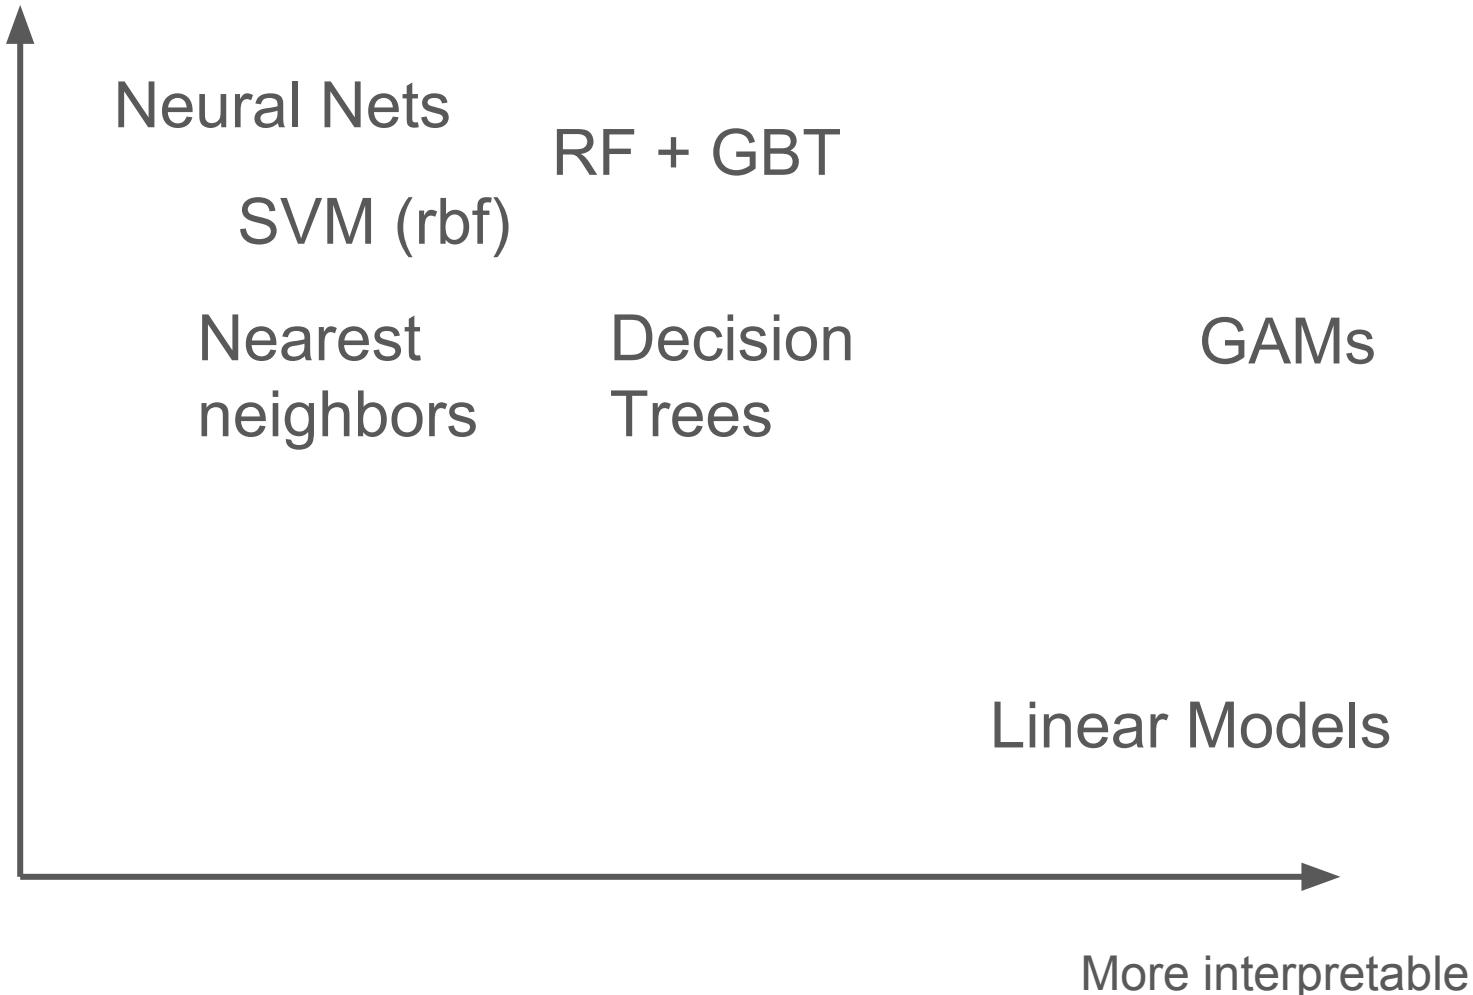
What are GAMs

but for us...

$$y = f_1(x_1) + f_2(x_2) + f_3(x_3) + c$$

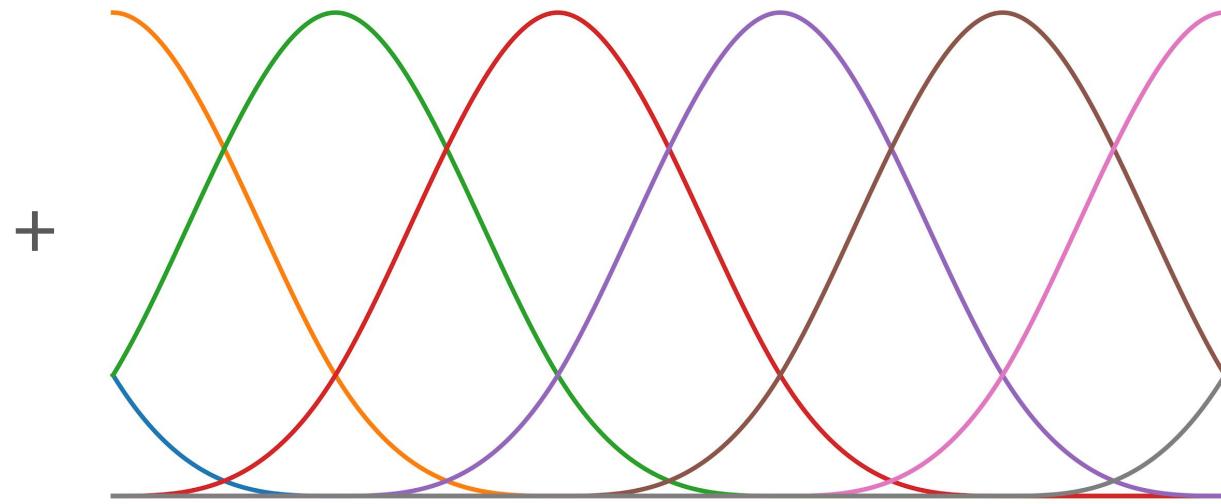
What are GAMs

More predictive power



What are GAMs

GAMs = GLMs



+ **Smoothing Penalty**

Plan

- What are GAMs
- A brief tour or splines
- Smoothing
- GAMs in Python

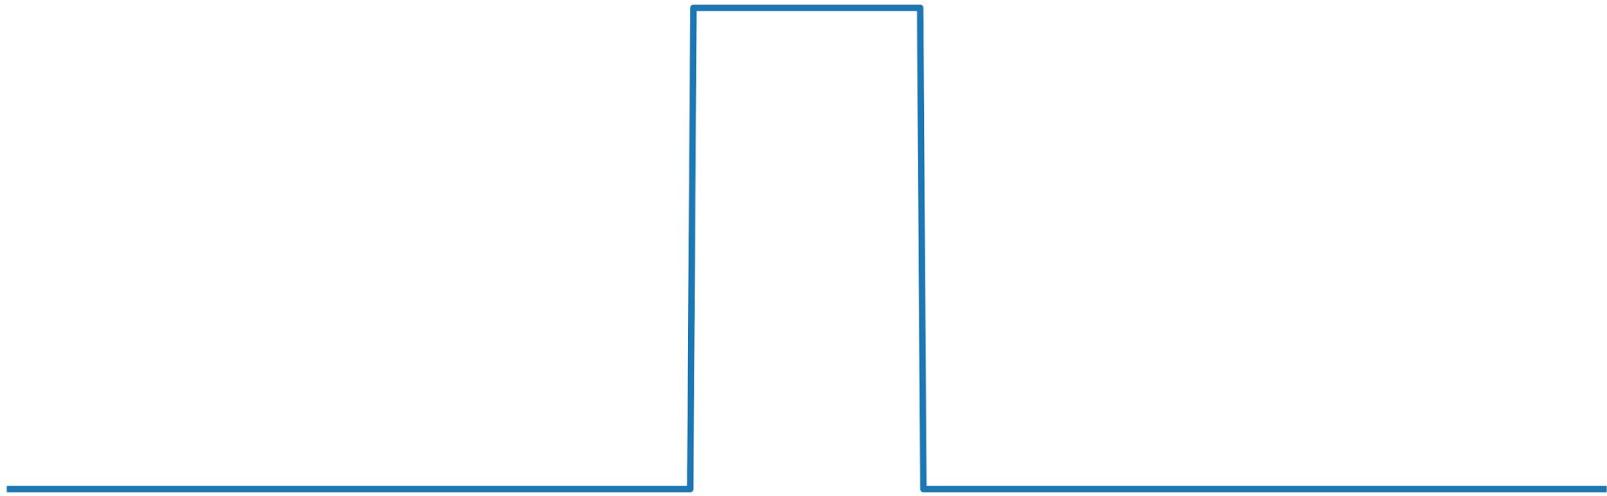
Splines

- Local support
- degree n spline has:
 - $n + 1$ segments of degree n
 - Continuous $(n-1)^{\text{th}}$ derivative

Splines

degree 0

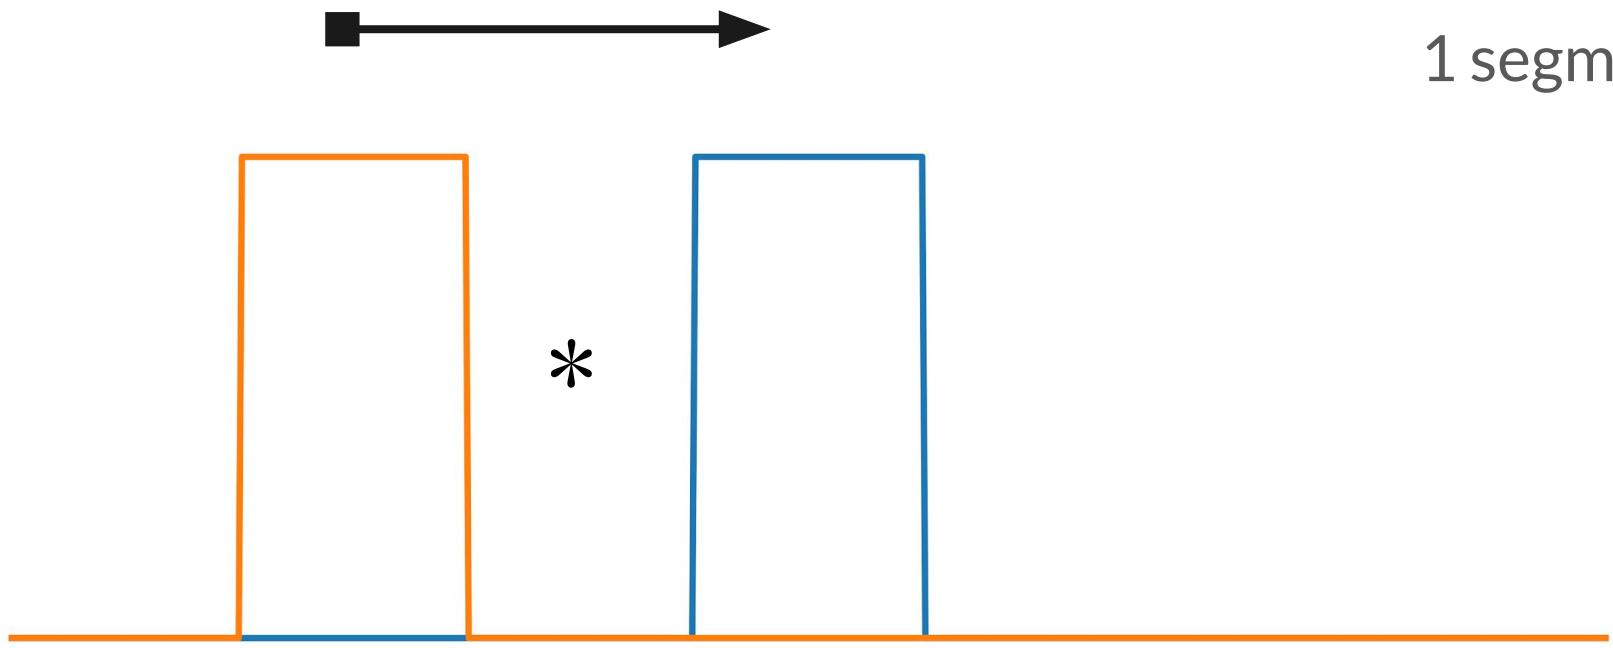
1 segment



Splines

degree 0

1 segment

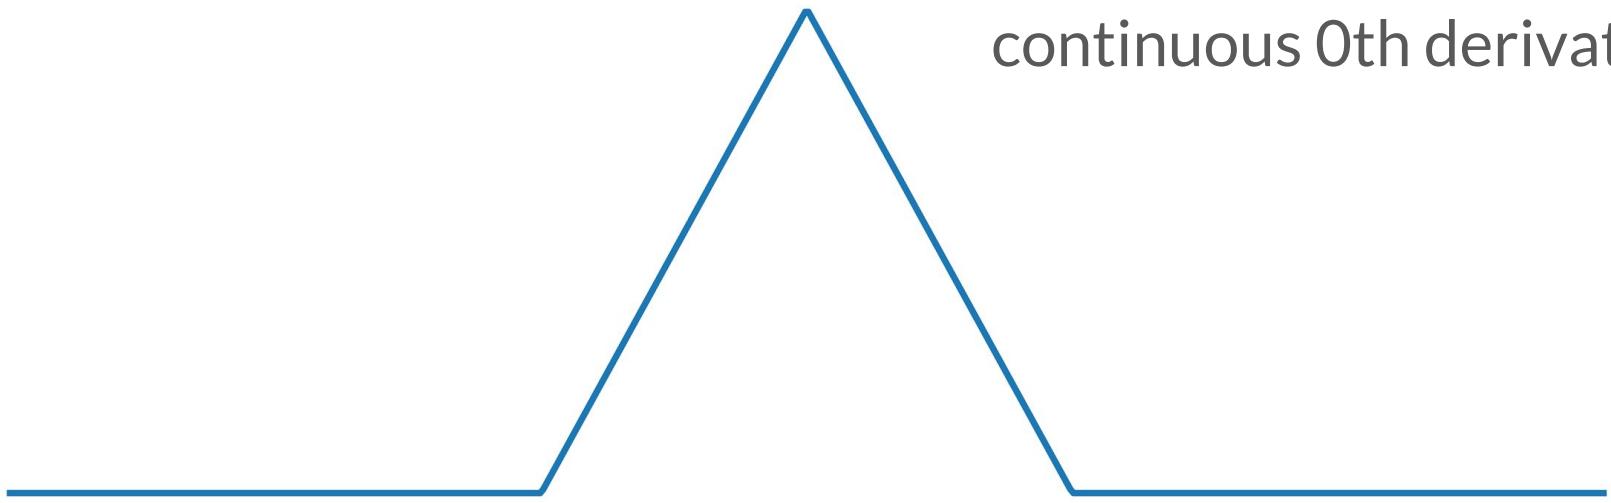


Splines

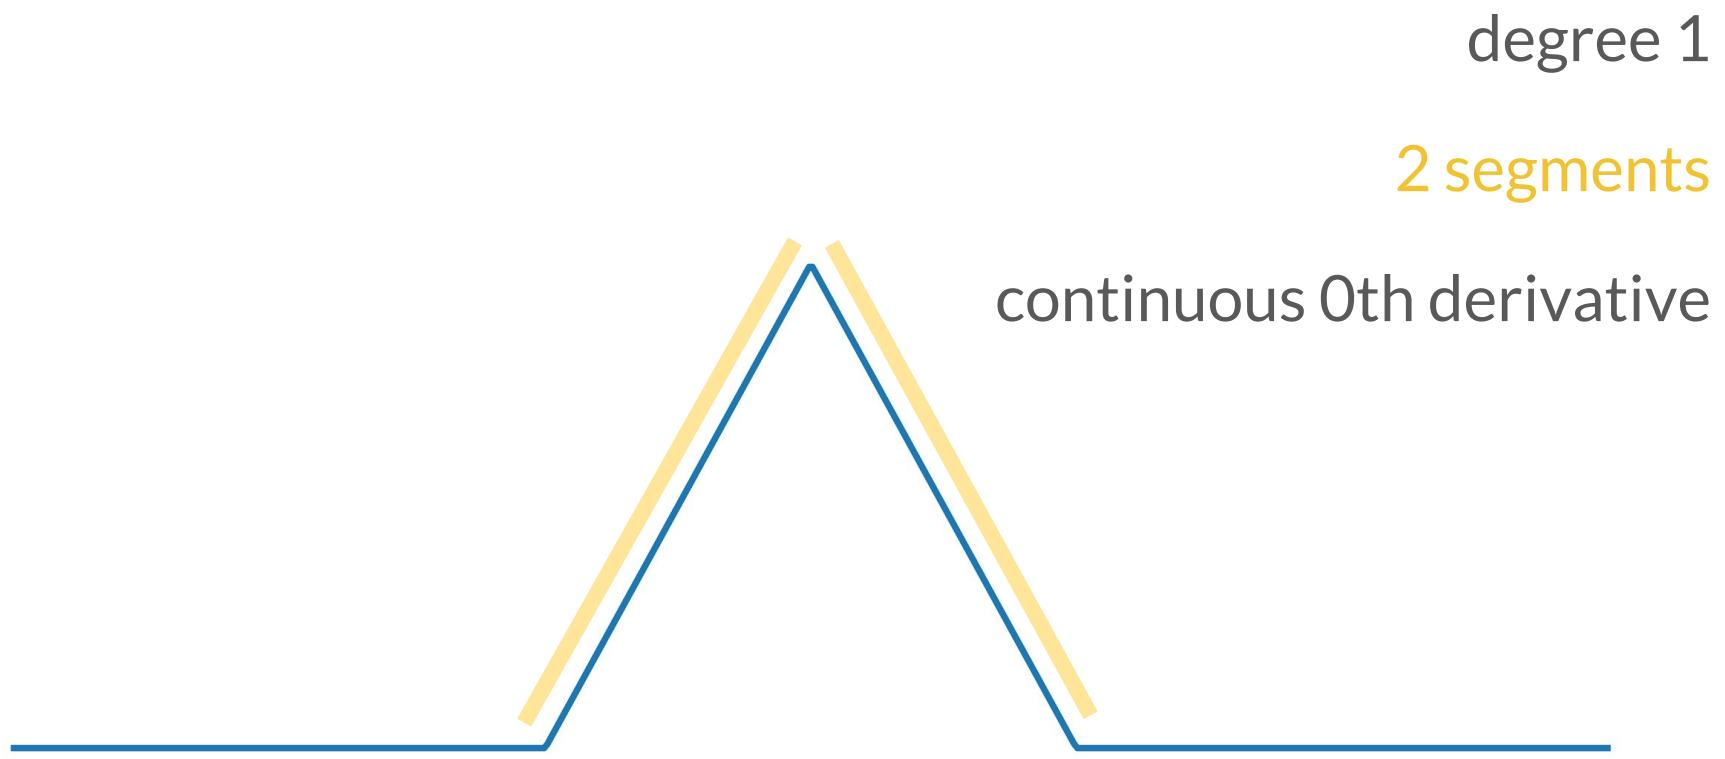
degree 1

2 segments

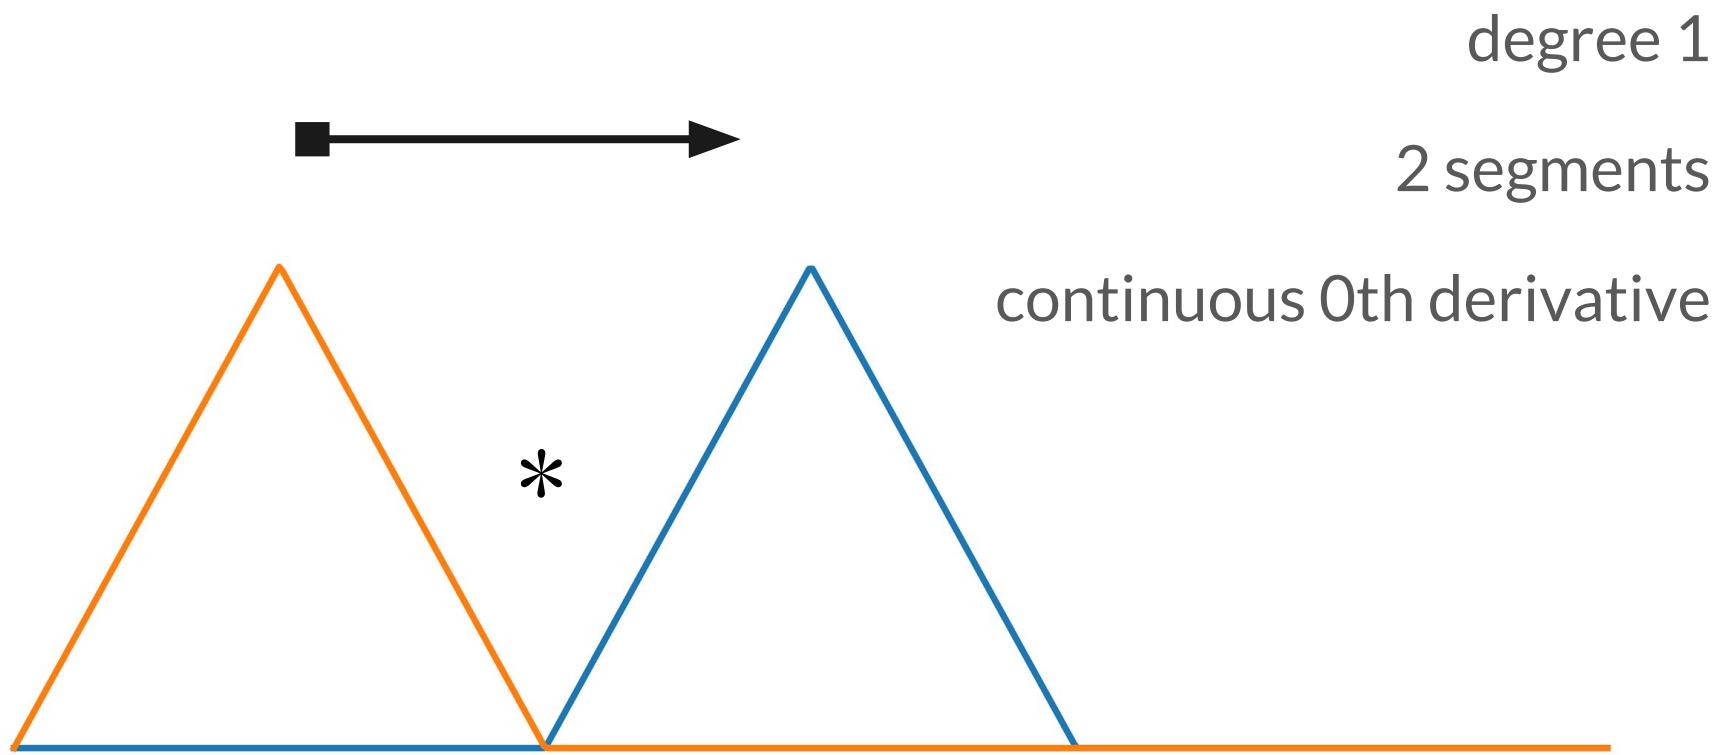
continuous 0th derivative



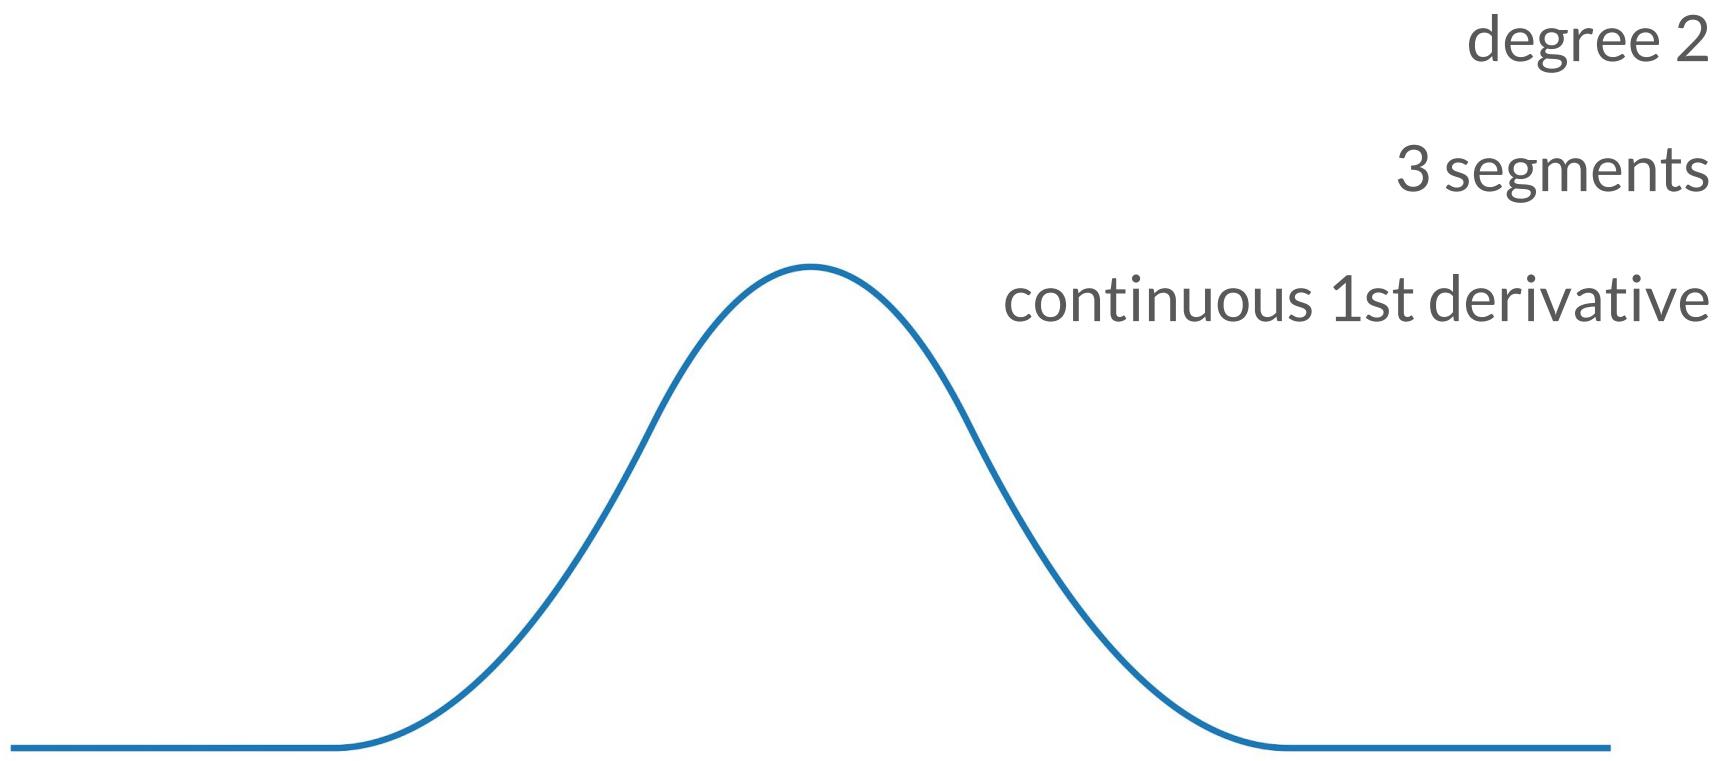
Splines



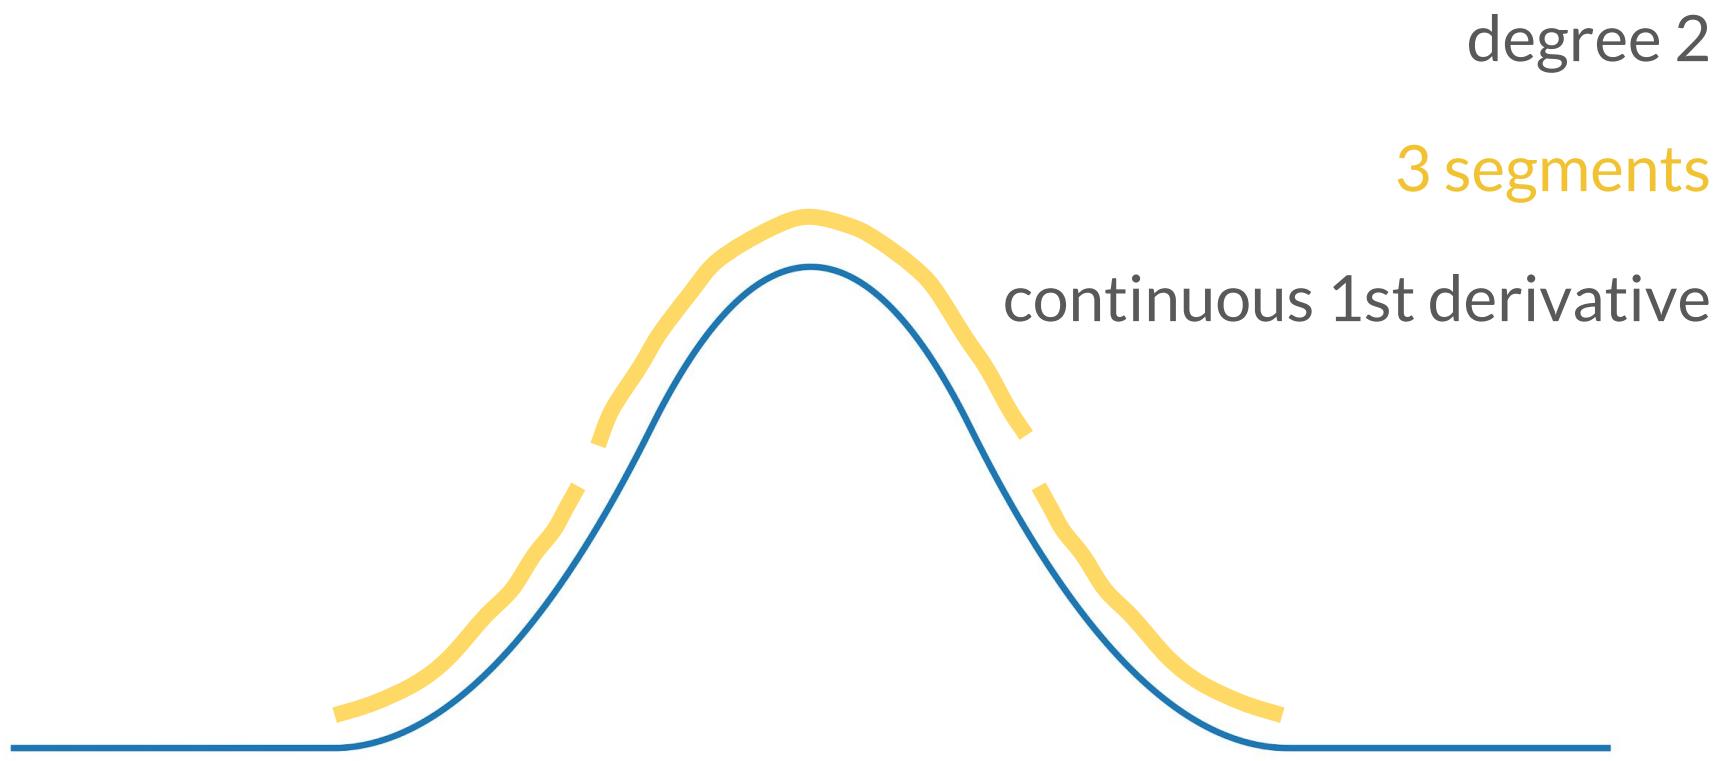
Splines



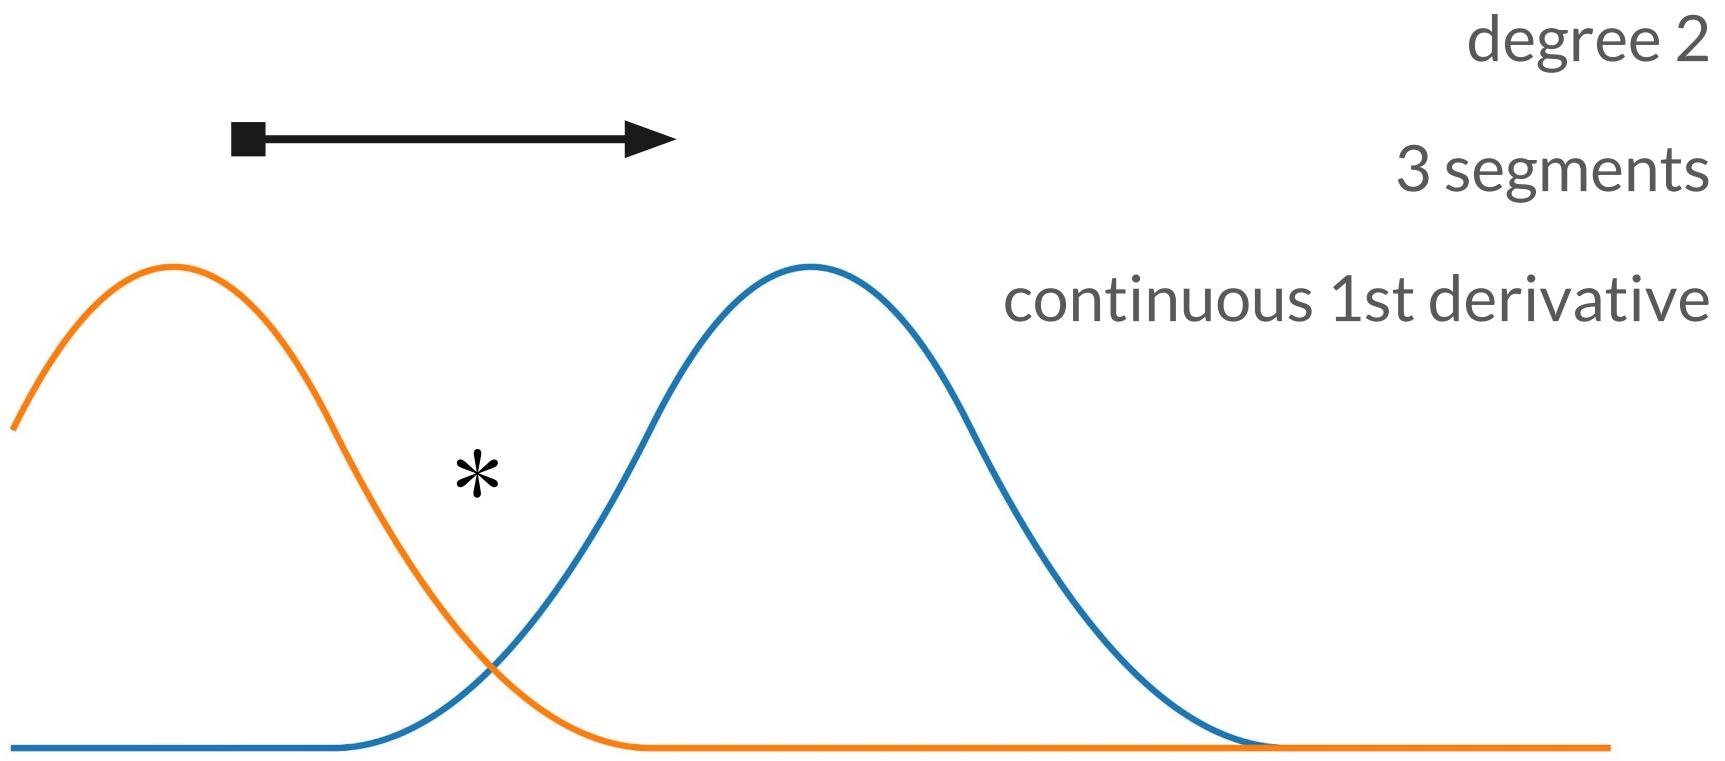
Splines



Splines



Splines

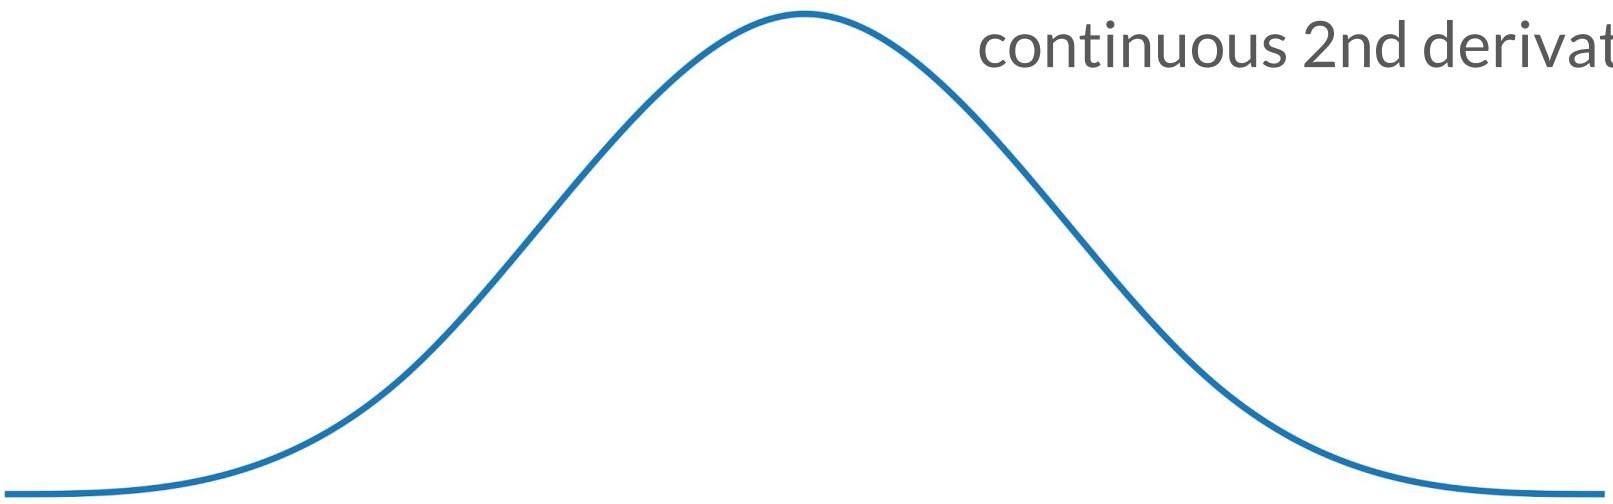


Splines

degree 3

4 segments

continuous 2nd derivative



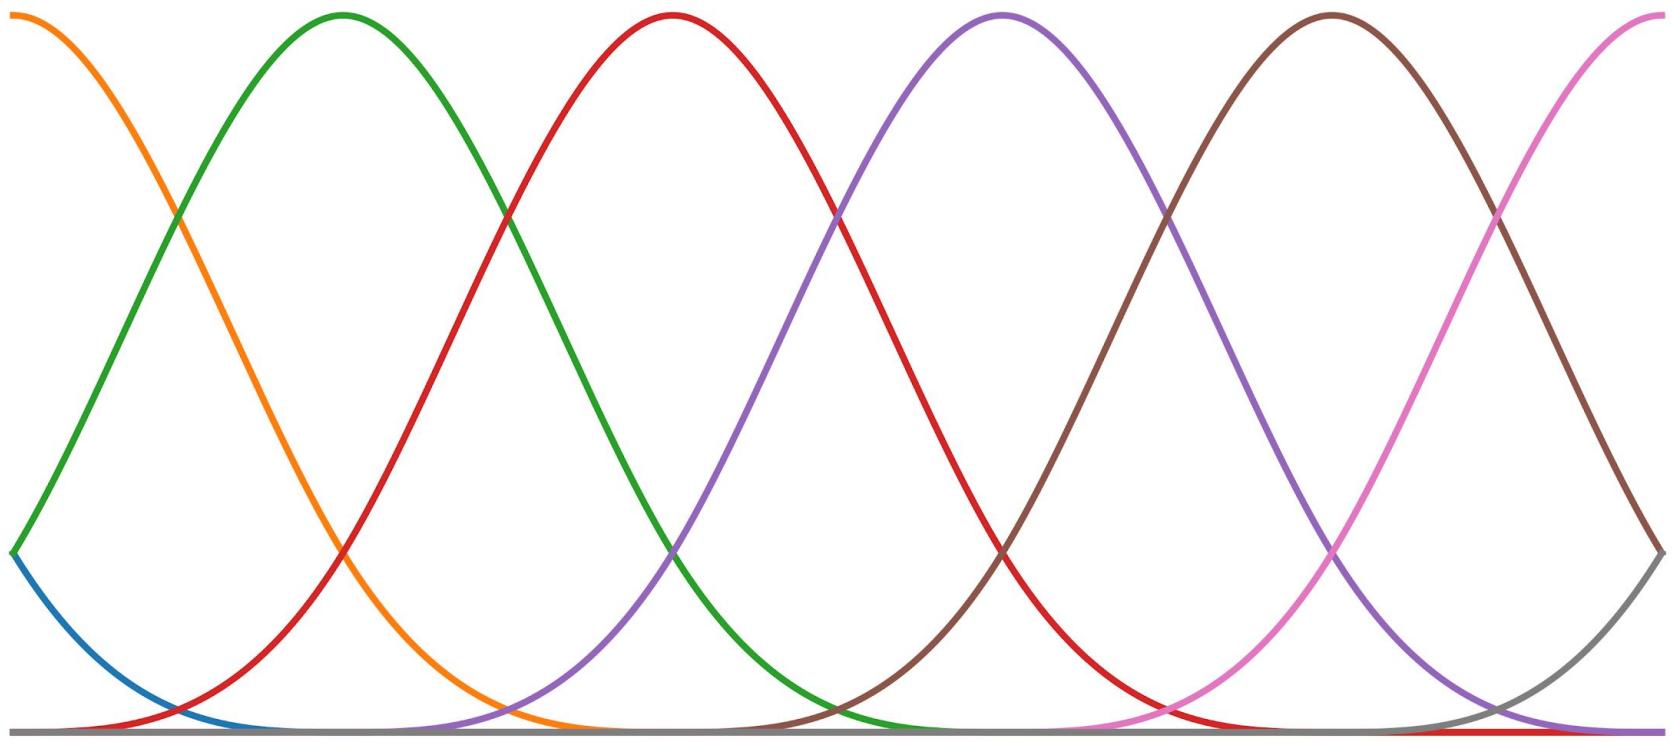
Splines

Spline -> Gaussian

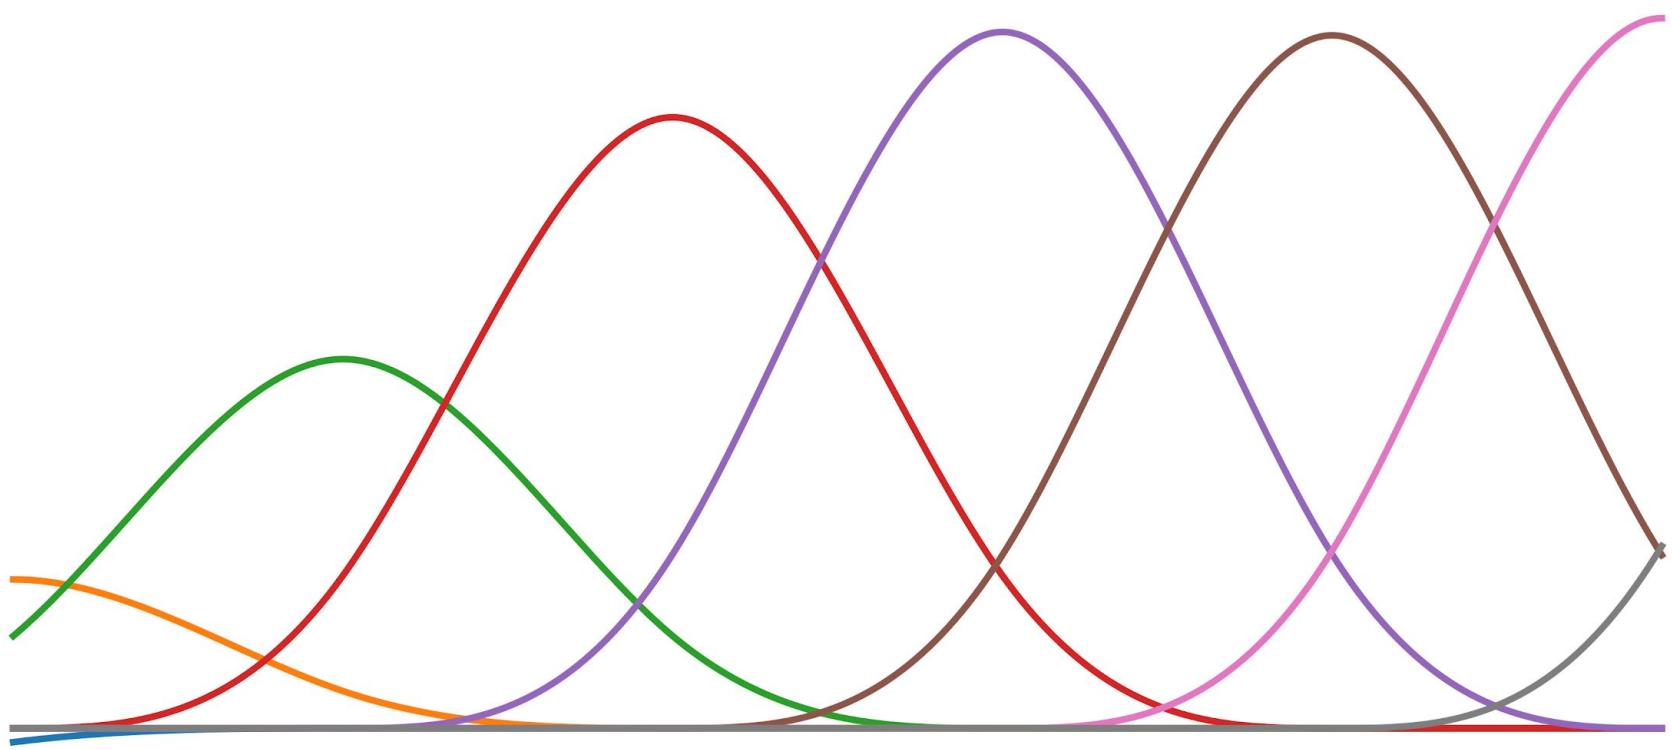
as

Spline Degree -> Infinity

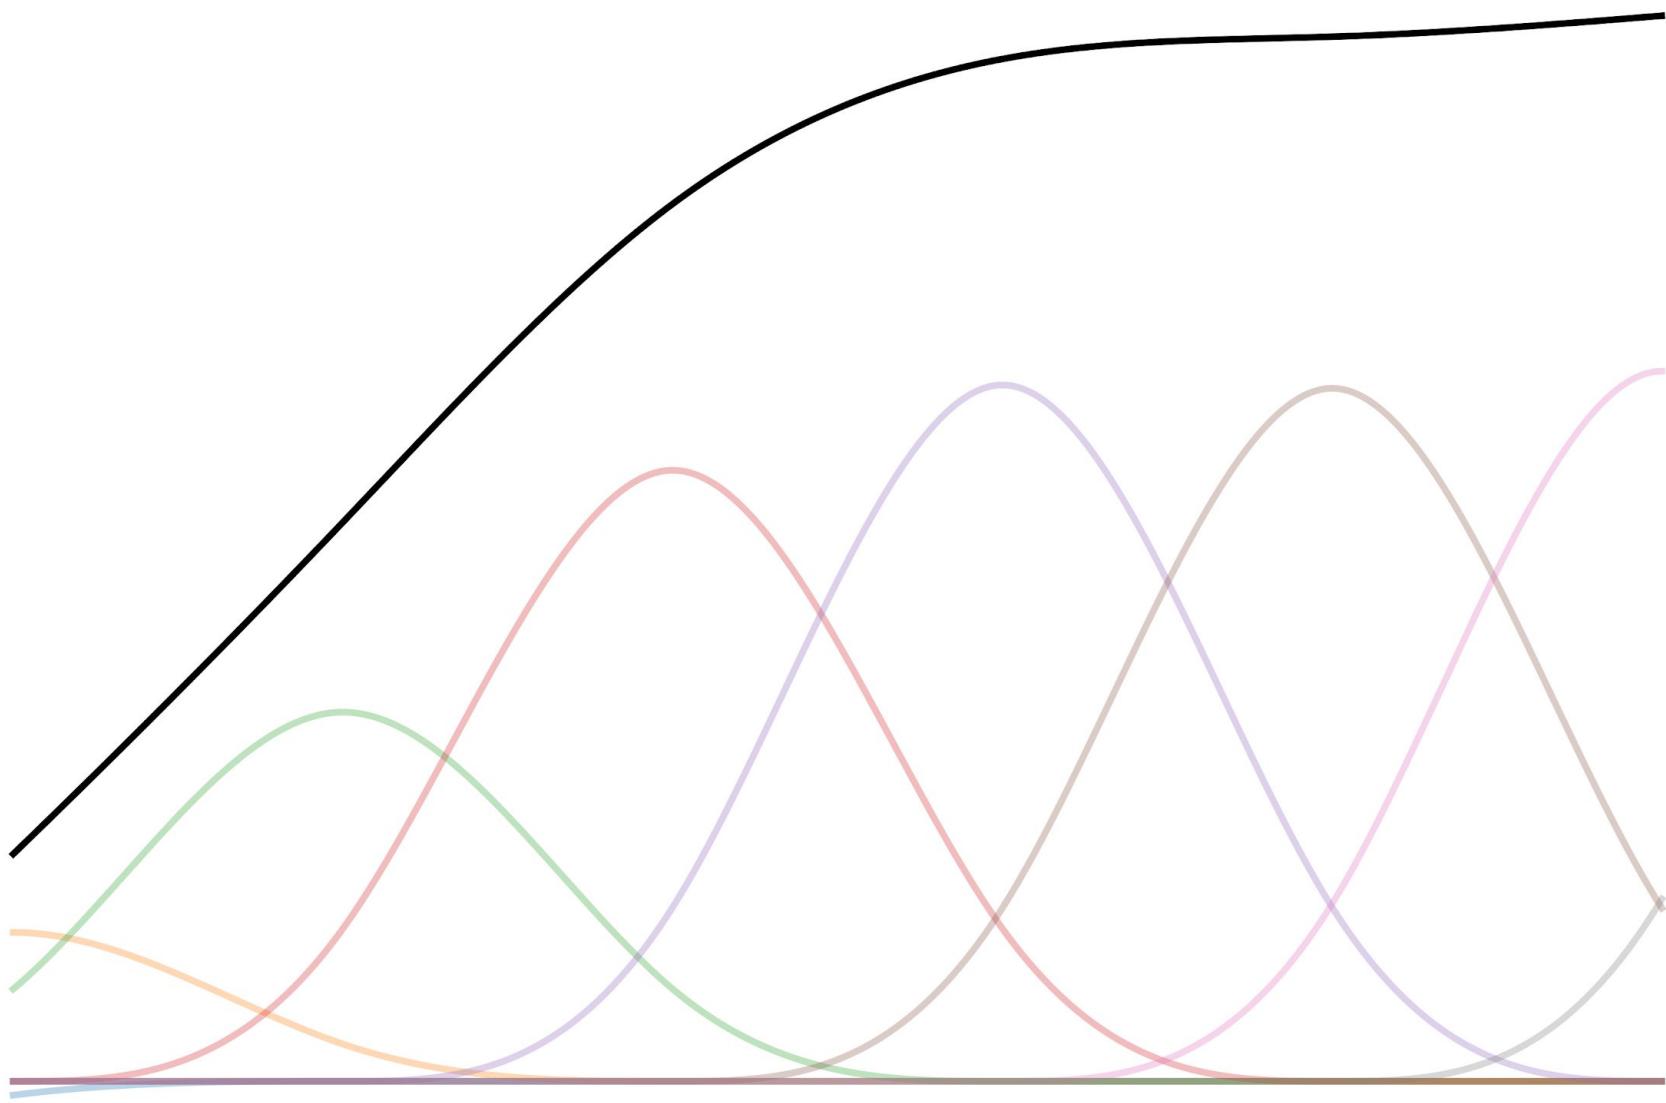
Splines



Splines



Splines



A simple linear model:

$$y_i = X_i \beta + \epsilon_i$$

$$\epsilon_i = \mathcal{N}(0, \sigma^2)$$

Some math...

A simple GAM:

$$y_i = f(x_i) + \epsilon_i$$

$$\epsilon_i = \mathcal{N}(0, \sigma^2)$$

Some math...

$$f(x) = \sum_{k=1}^q b_k(x) \beta_k$$

Some math...

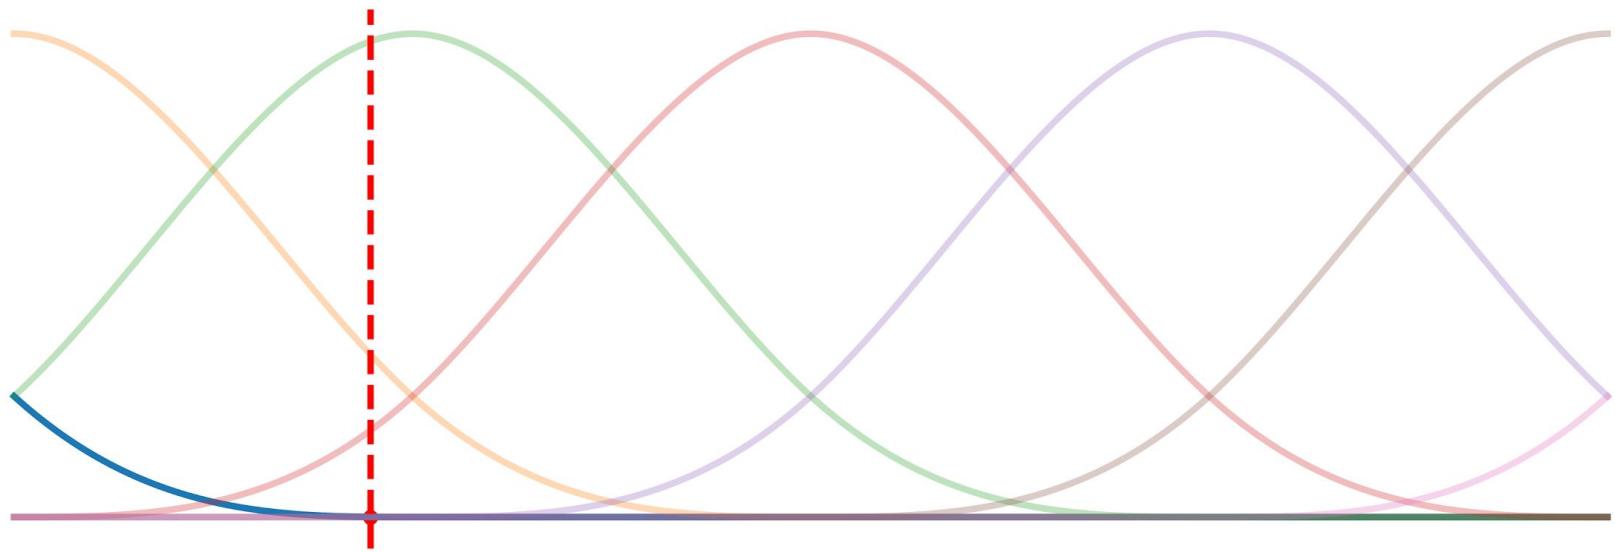
$$\sum_{k=1}^q b_k(x) \beta_k$$

Spline functions

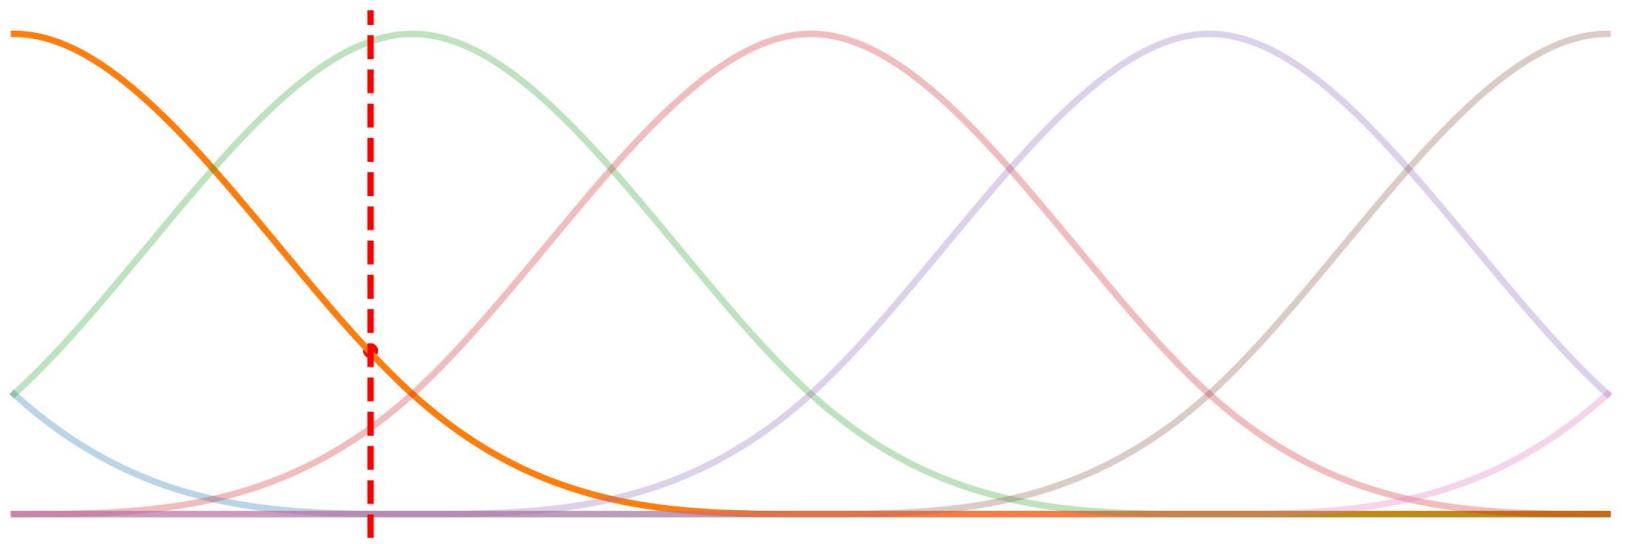
coefficient

The diagram illustrates a mathematical expression involving a sum of spline basis functions. The expression is $\sum_{k=1}^q b_k(x) \beta_k$. A curved arrow originates from the text "Spline functions" and points to the term $b_k(x)$. Another curved arrow originates from the text "coefficient" and points to the term β_k .

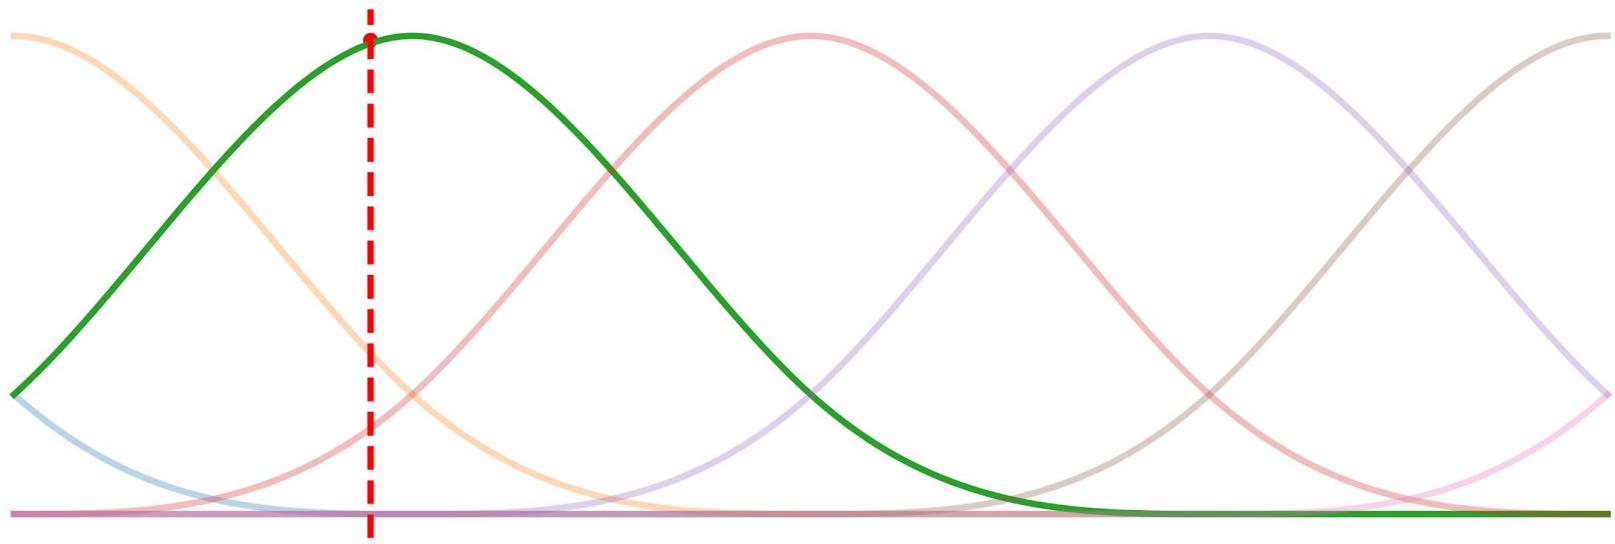
Some math...



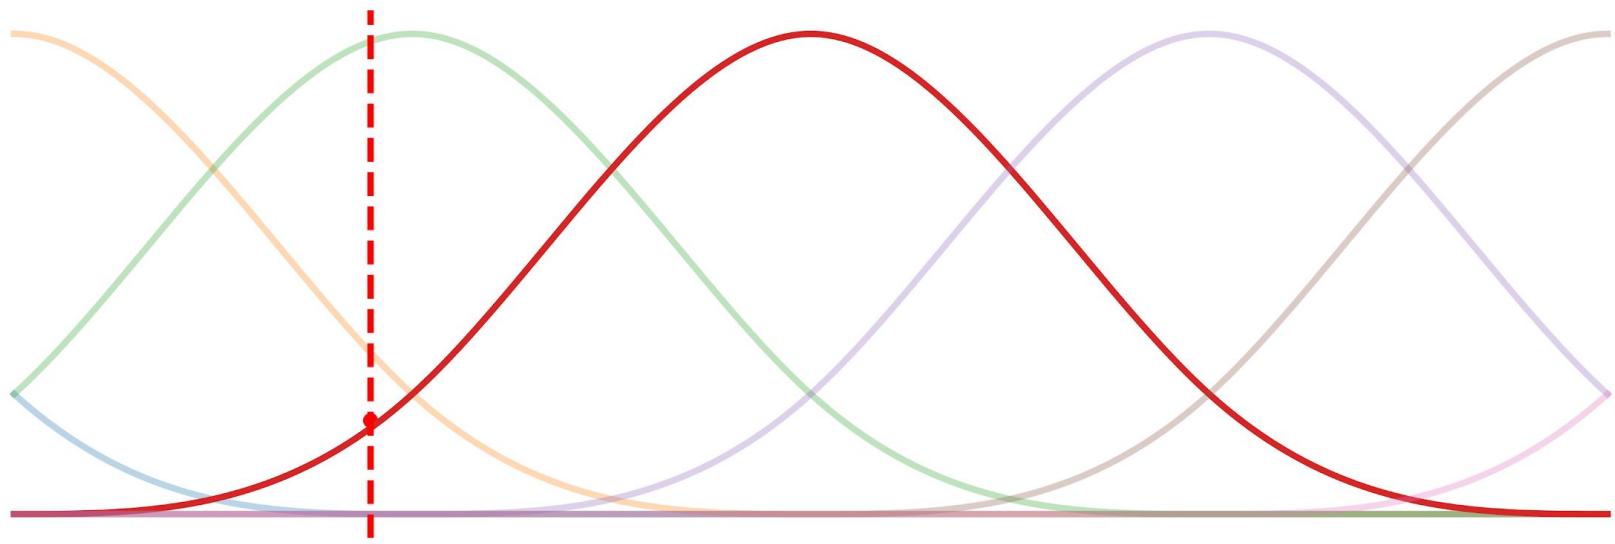
Some math...



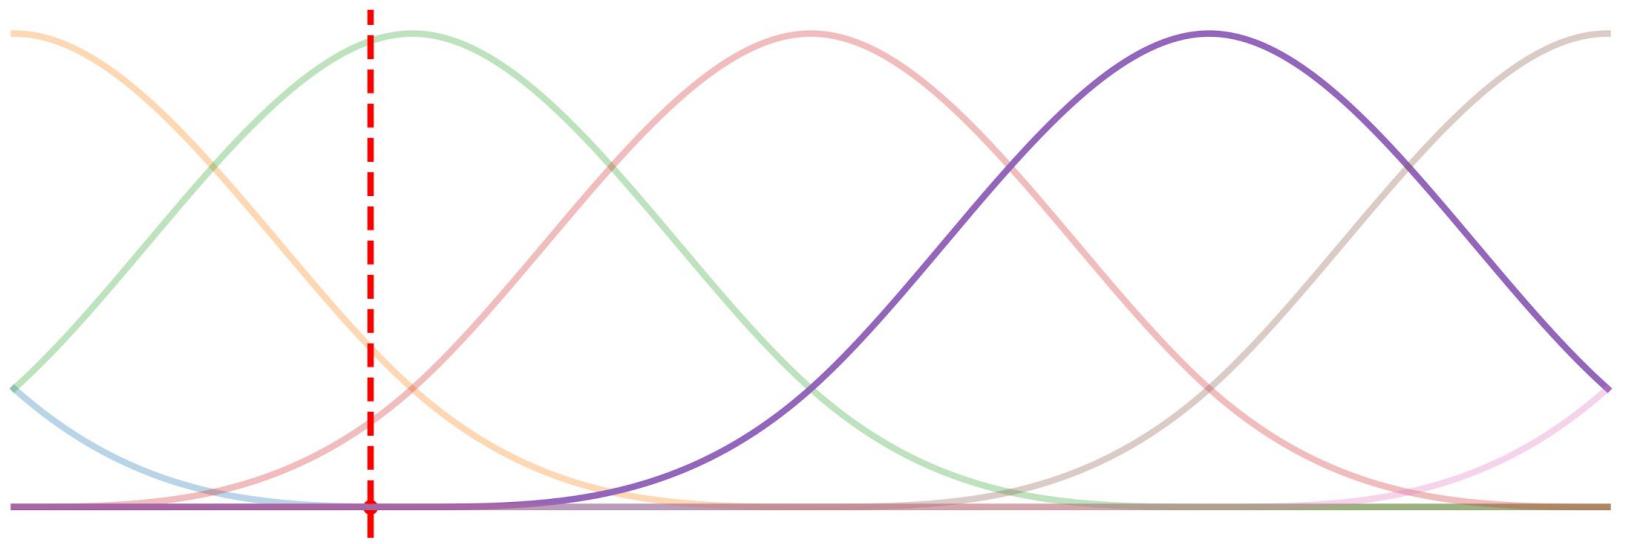
Some math...



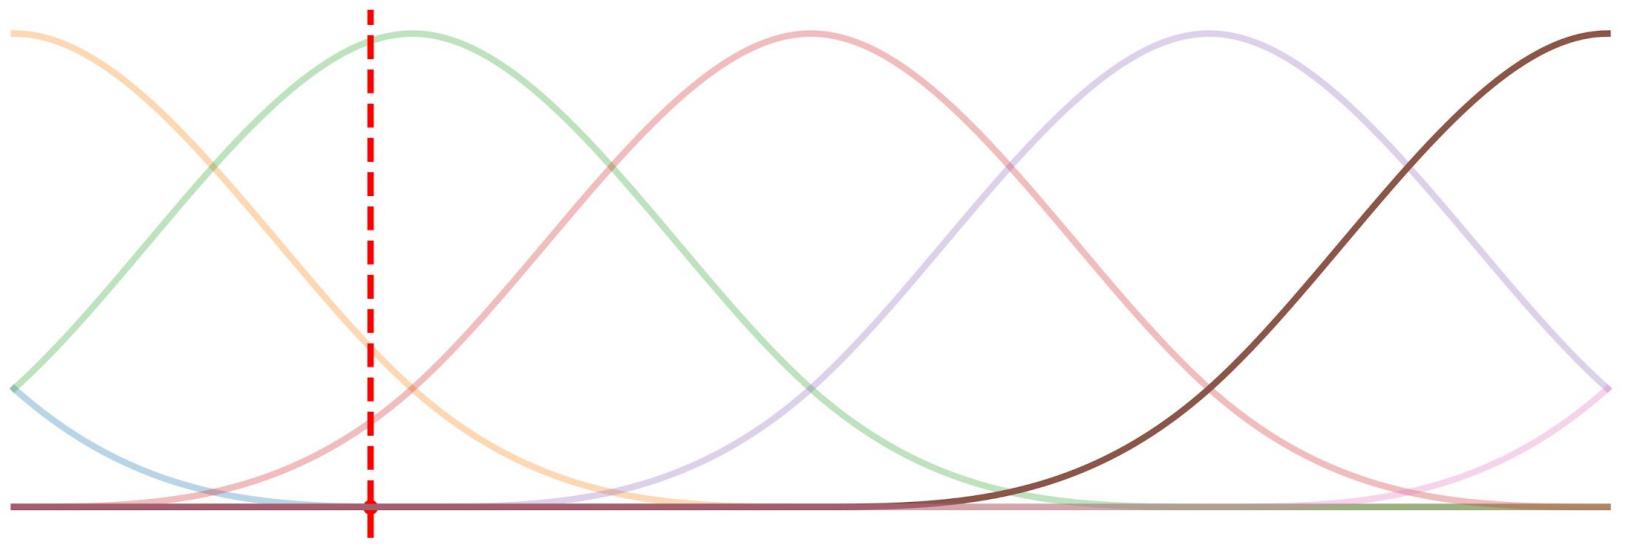
Some math...



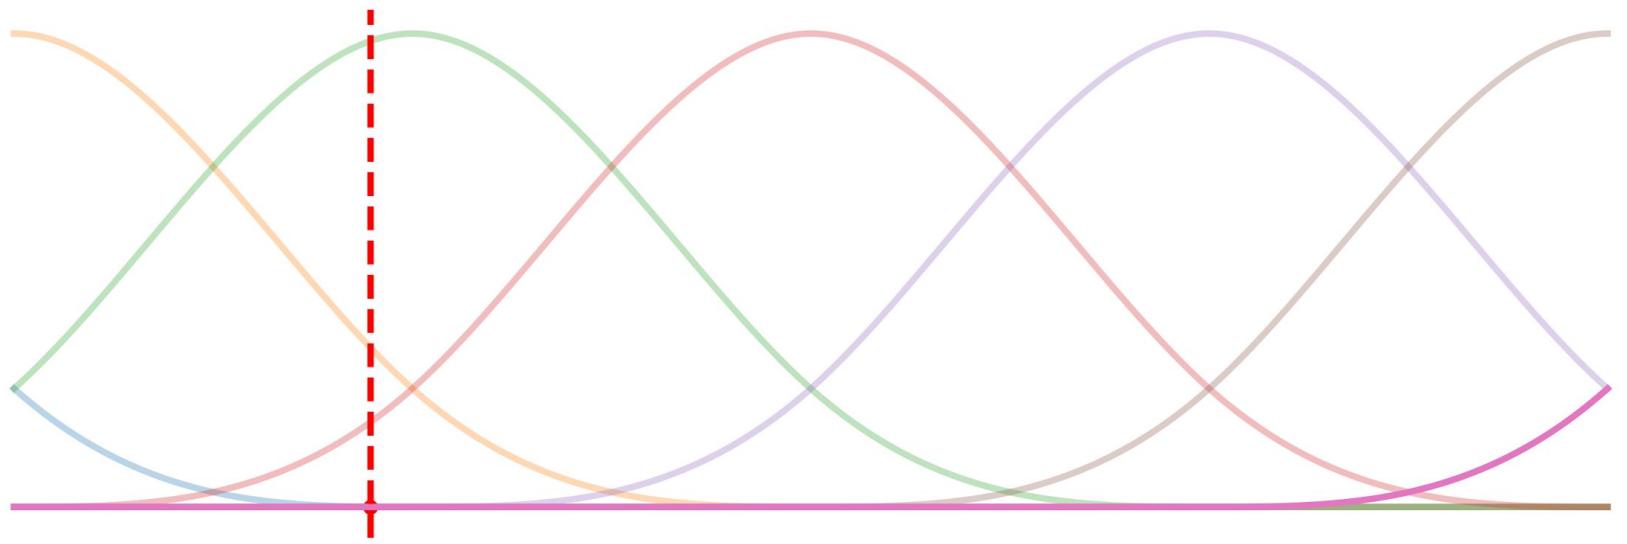
Some math...



Some math...



Some math...



Some math...

$$\sum_{k=1}^q b_k(x) \beta_k$$

Some math...

$$\sum_{k=1}^q \hat{x}_k \beta_k$$

Some math...

$$\hat{X}^\beta$$

Some math...

$$y = \hat{X}\beta + \epsilon$$

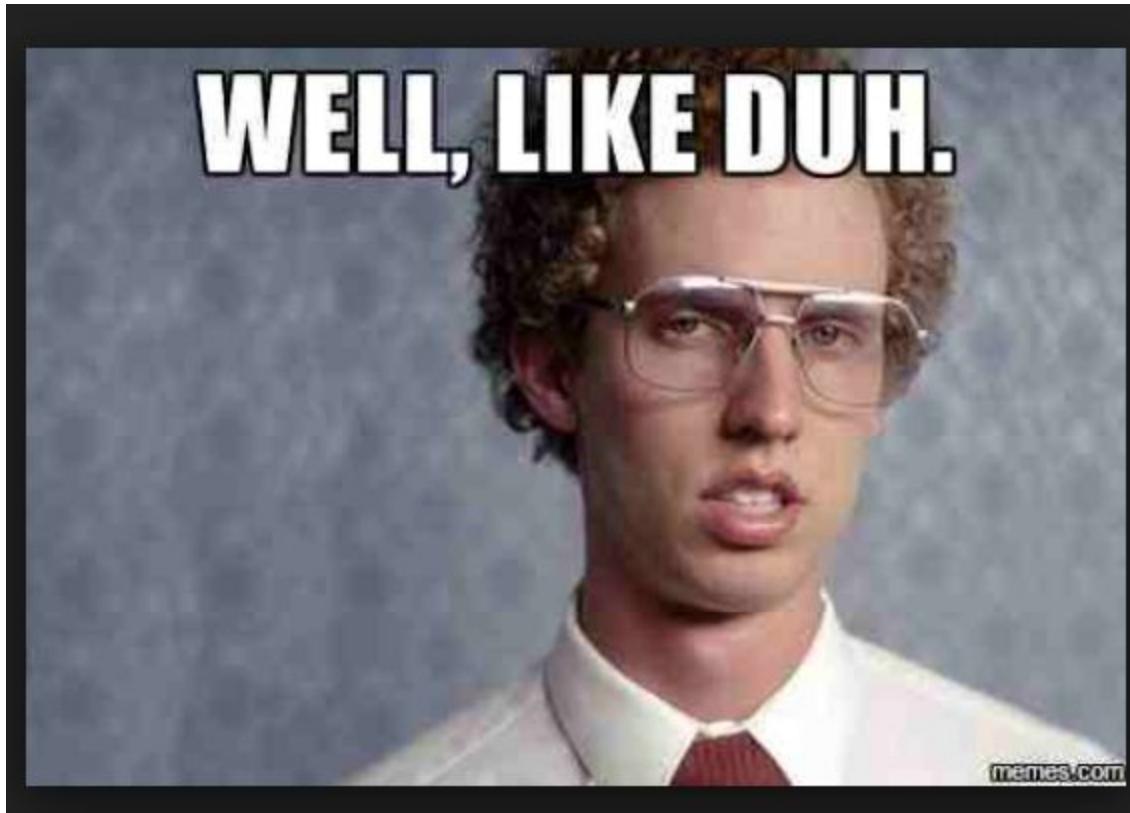
Some math...

$$y_i = \hat{X}_i \beta + \epsilon_i$$

Some math...

GAM = big GLM

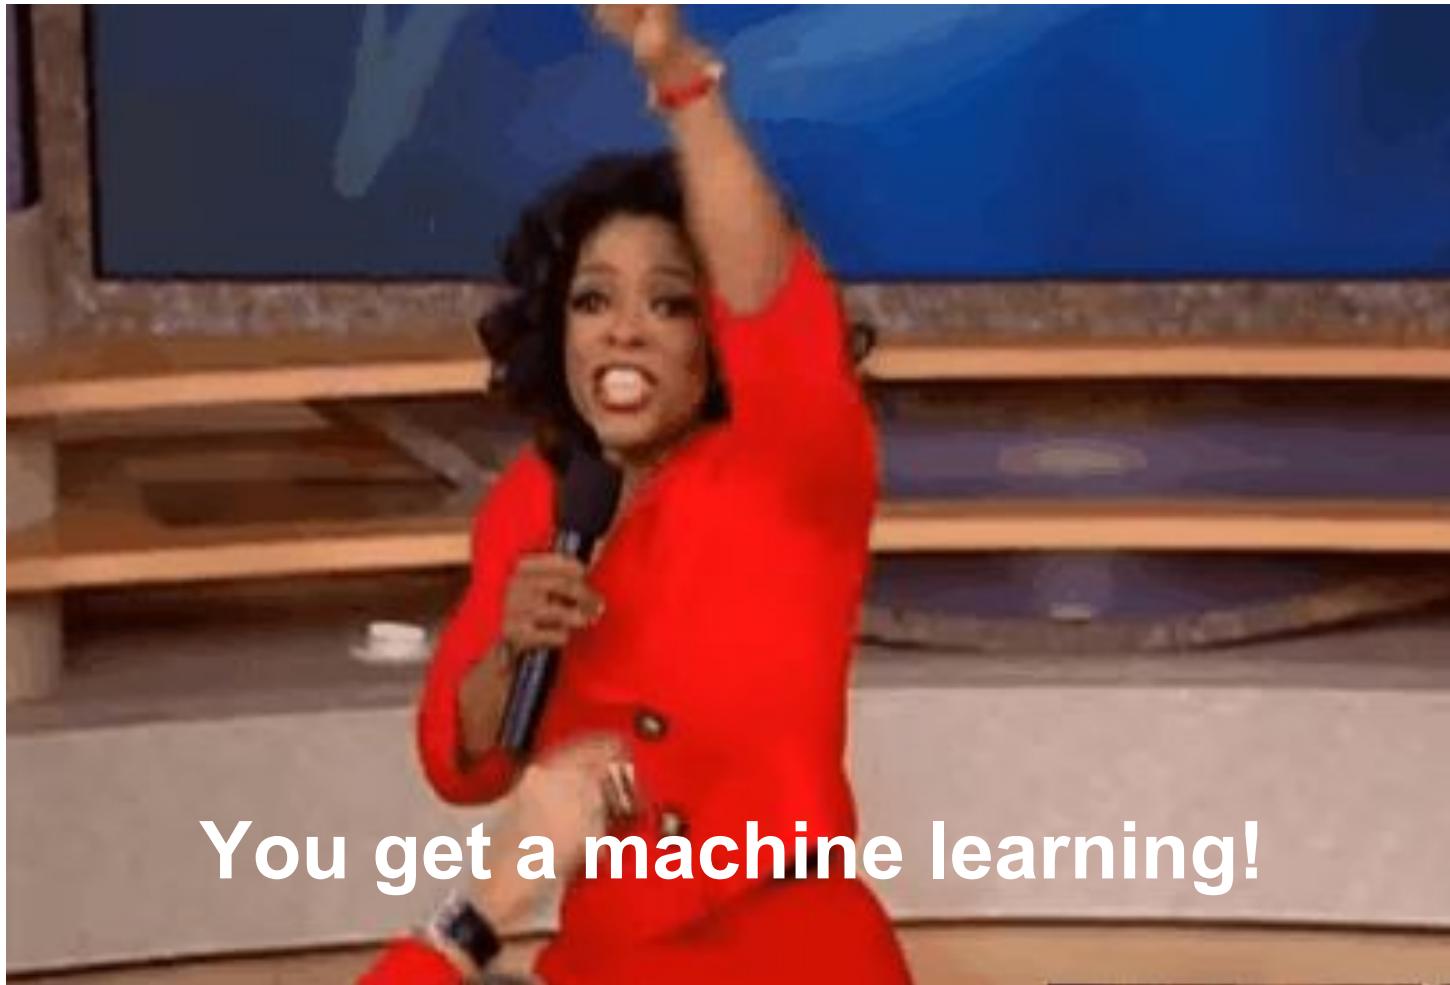
Some math...



Lots of **GLM** theory for free

- Confidence intervals
- Prediction intervals
- non-Normal observations
- Feature selection with p-values
- Fast cross-validation via GCV

Free stuff!



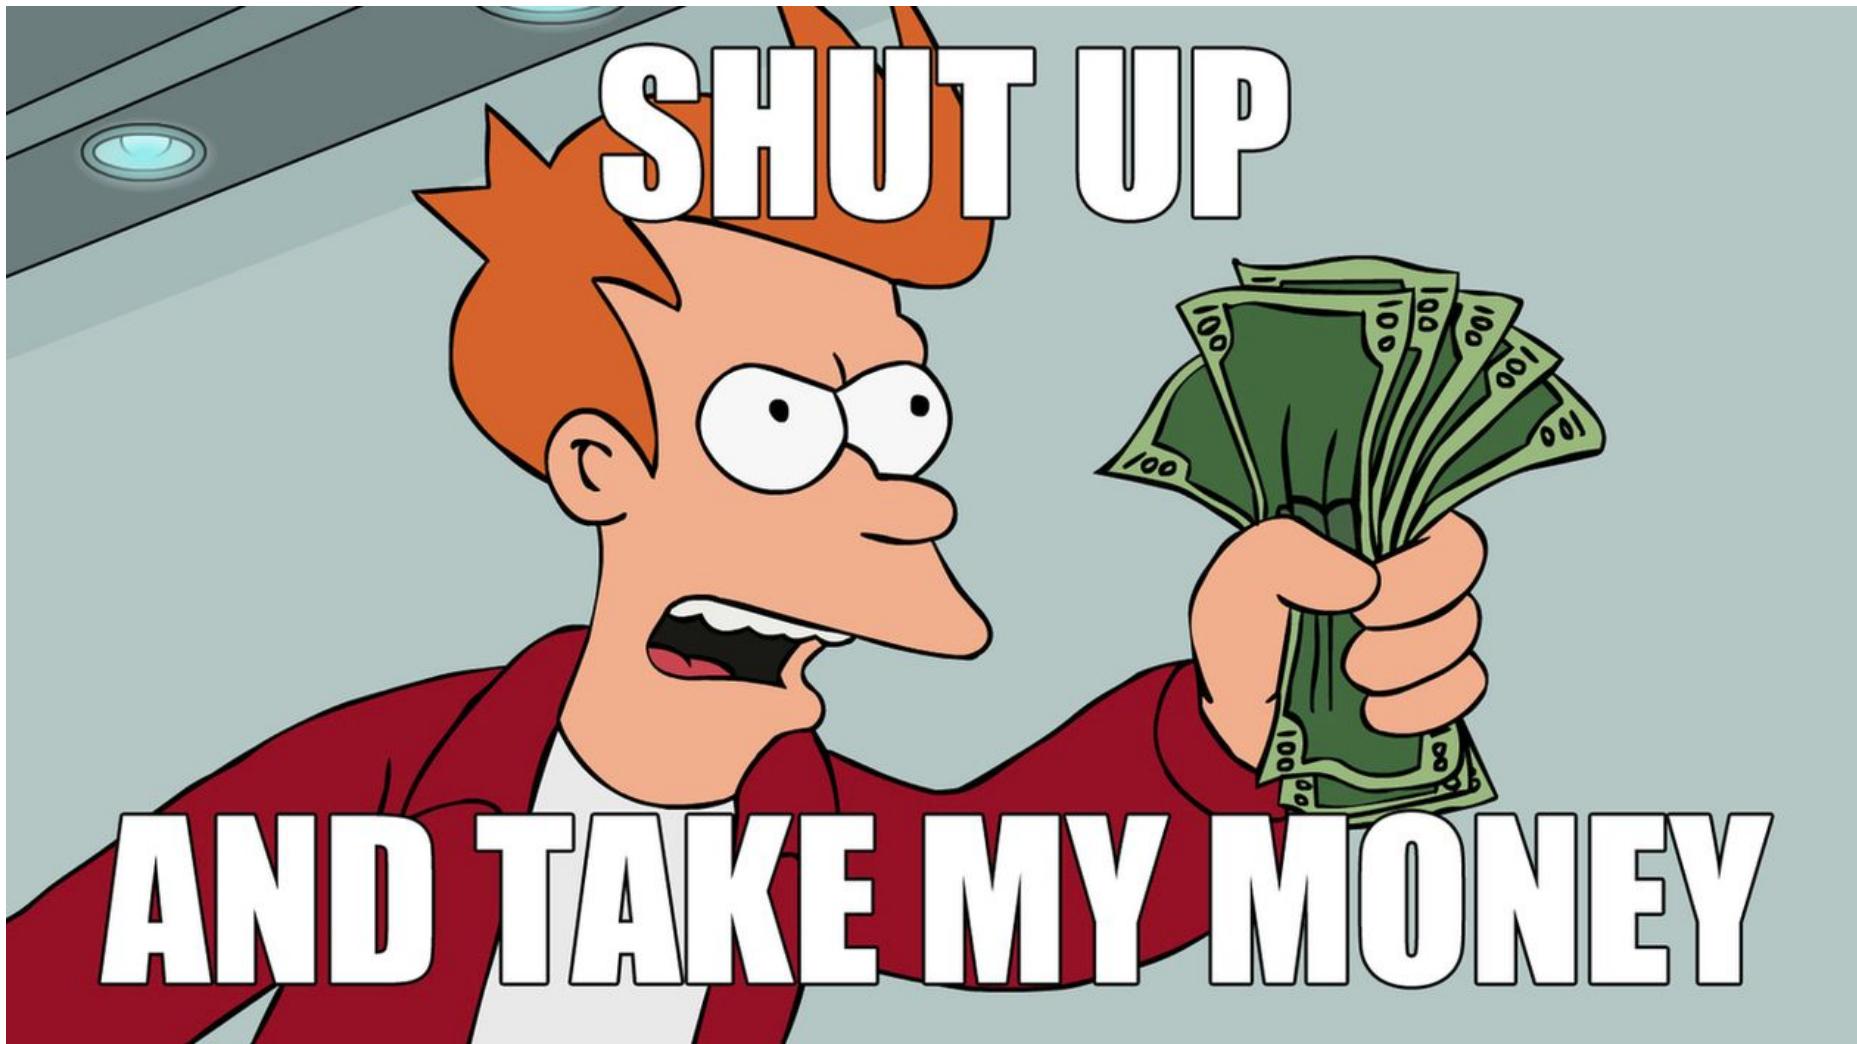
You get a machine learning!

Free stuff!

And specific GAM stuff

- Intuitive regularization via smoothing
- Automatically model non-linearities
- Constraints (convexity, monotonicity, periodicity, ...)
- Controlled extrapolation

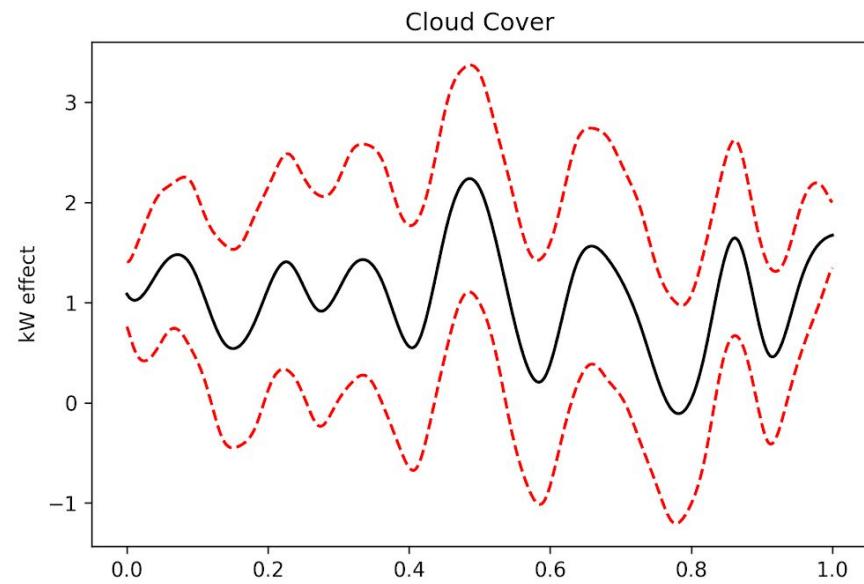
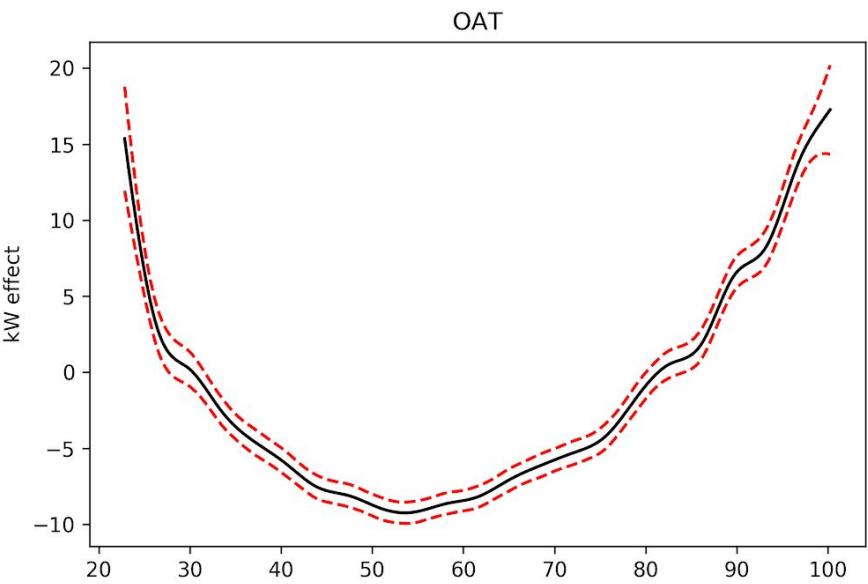
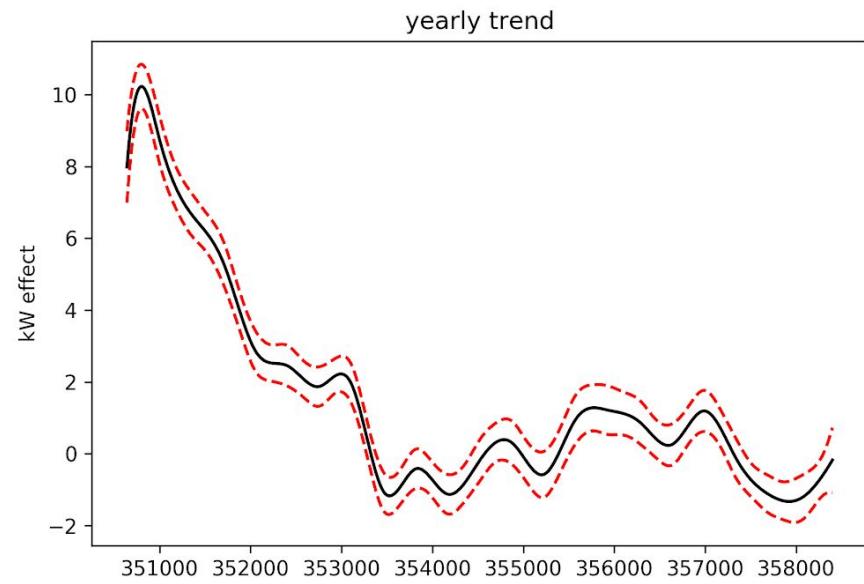
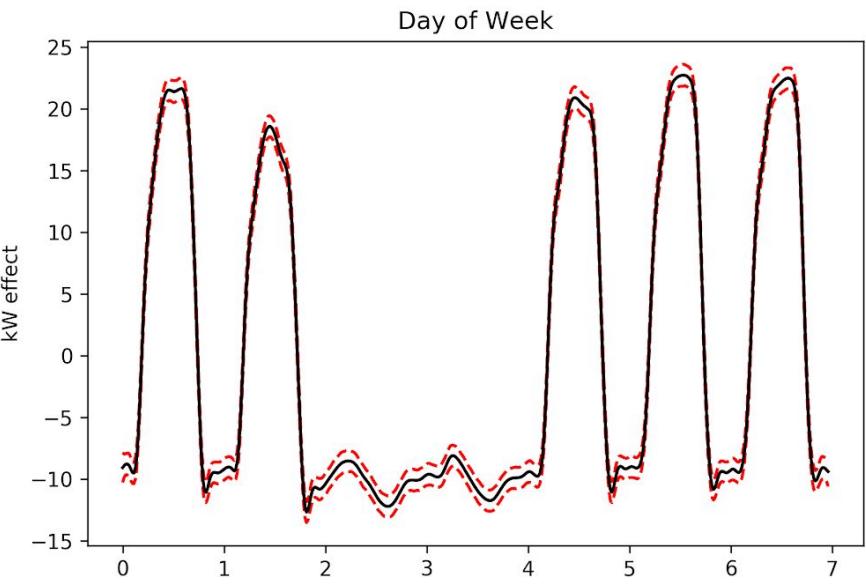
Free stuff!



Plan

- What are GAMs
- A brief tour or splines
- Smoothing

Smoothing



Smoothing

Penalize excessive wigginess

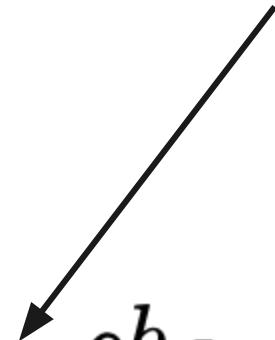
$$\text{wigginess} = \int_a^b [f''(x)]^2 dx$$

Smoothing

minimize

$$\|y - X\beta\|^2 + \lambda \int_a^b [f''(x)]^2 dx$$

strength of smoothing

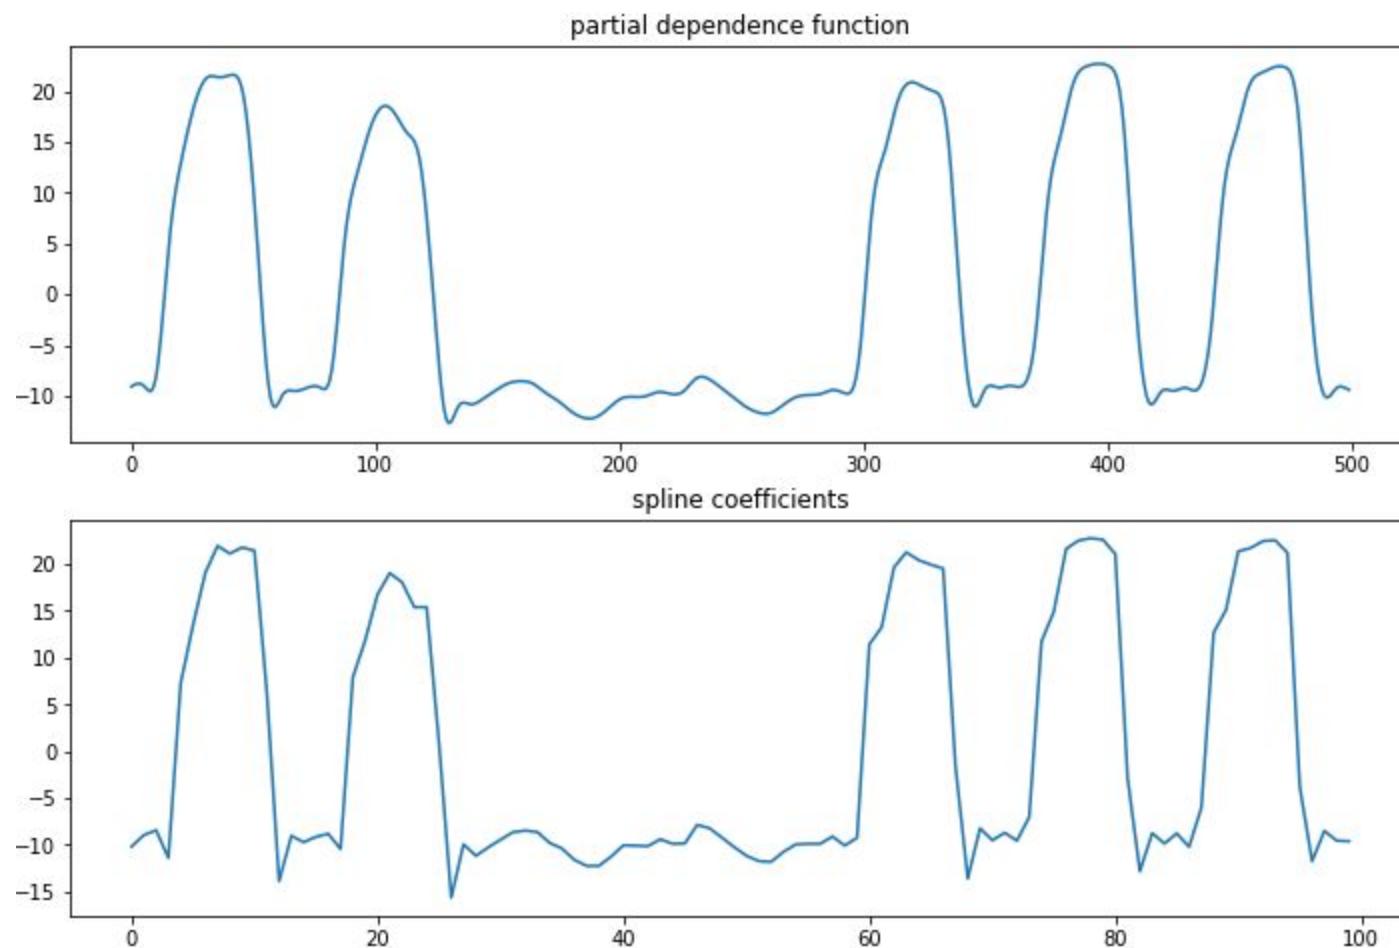


Smoothing

2 modifications:

- Compute on spline coefficients β
- Use numerical derivatives + integrals

Smoothing



Smoothing

minimize

$$\|y - X\beta\|^2 + \lambda\beta^\top S\beta$$

numerical derivatives

model coefficients

Smoothing

minimize

$$\beta = (X^\top X + \lambda S)^{-1} X^\top y$$

Smoothing

Similar to ridge regression

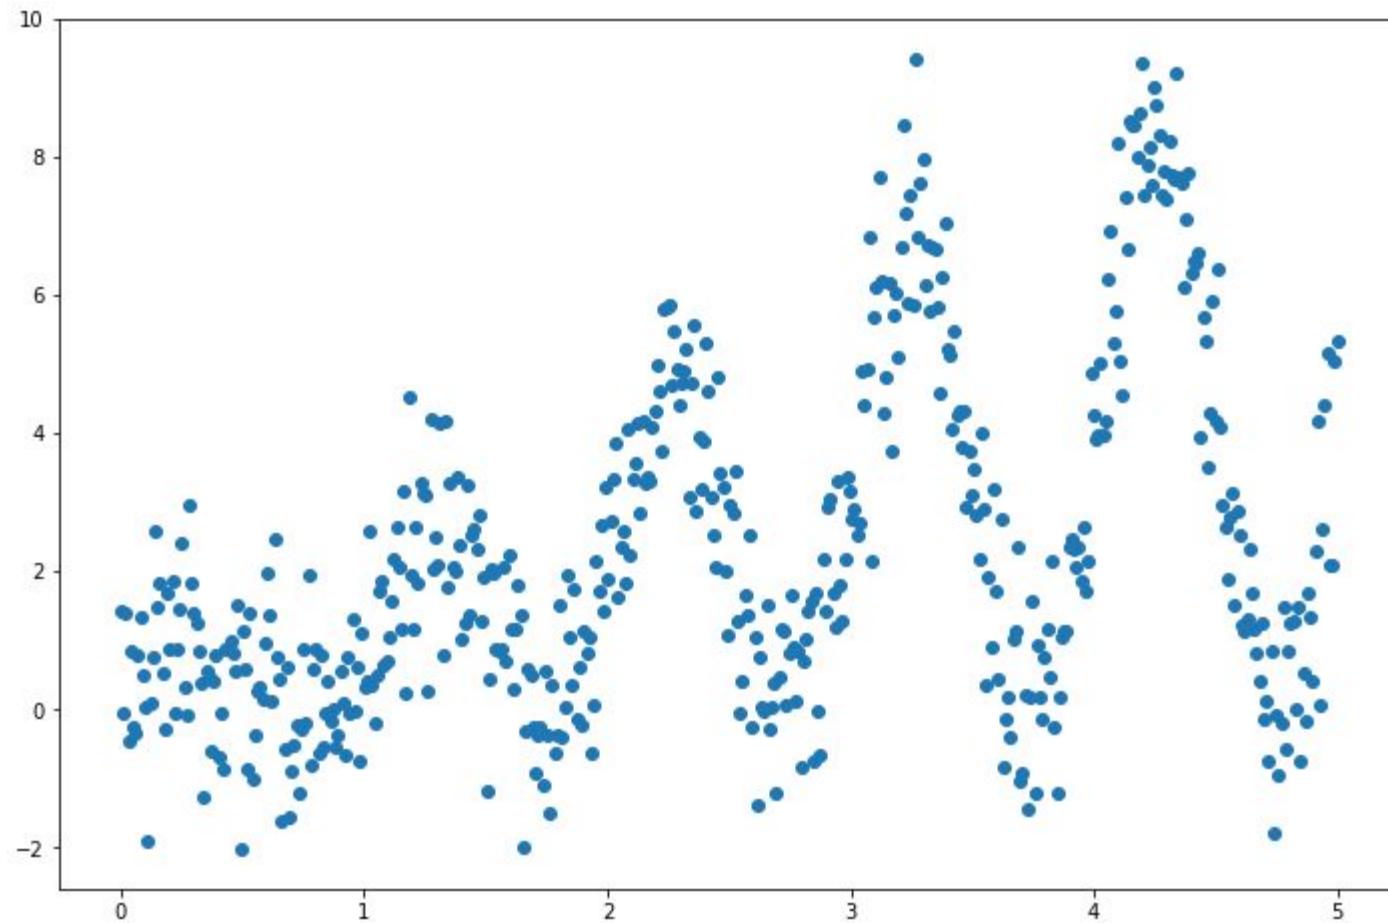
$$\beta = (X^\top X + \lambda I)^{-1} X^\top y$$

Smoothing

```
X = np.linspace(0, 5, 500)
y = np.sin(X * 2 * np.pi)*X + X + np.random.randn(len(X))

plt.scatter(X, y)
```

Smoothing



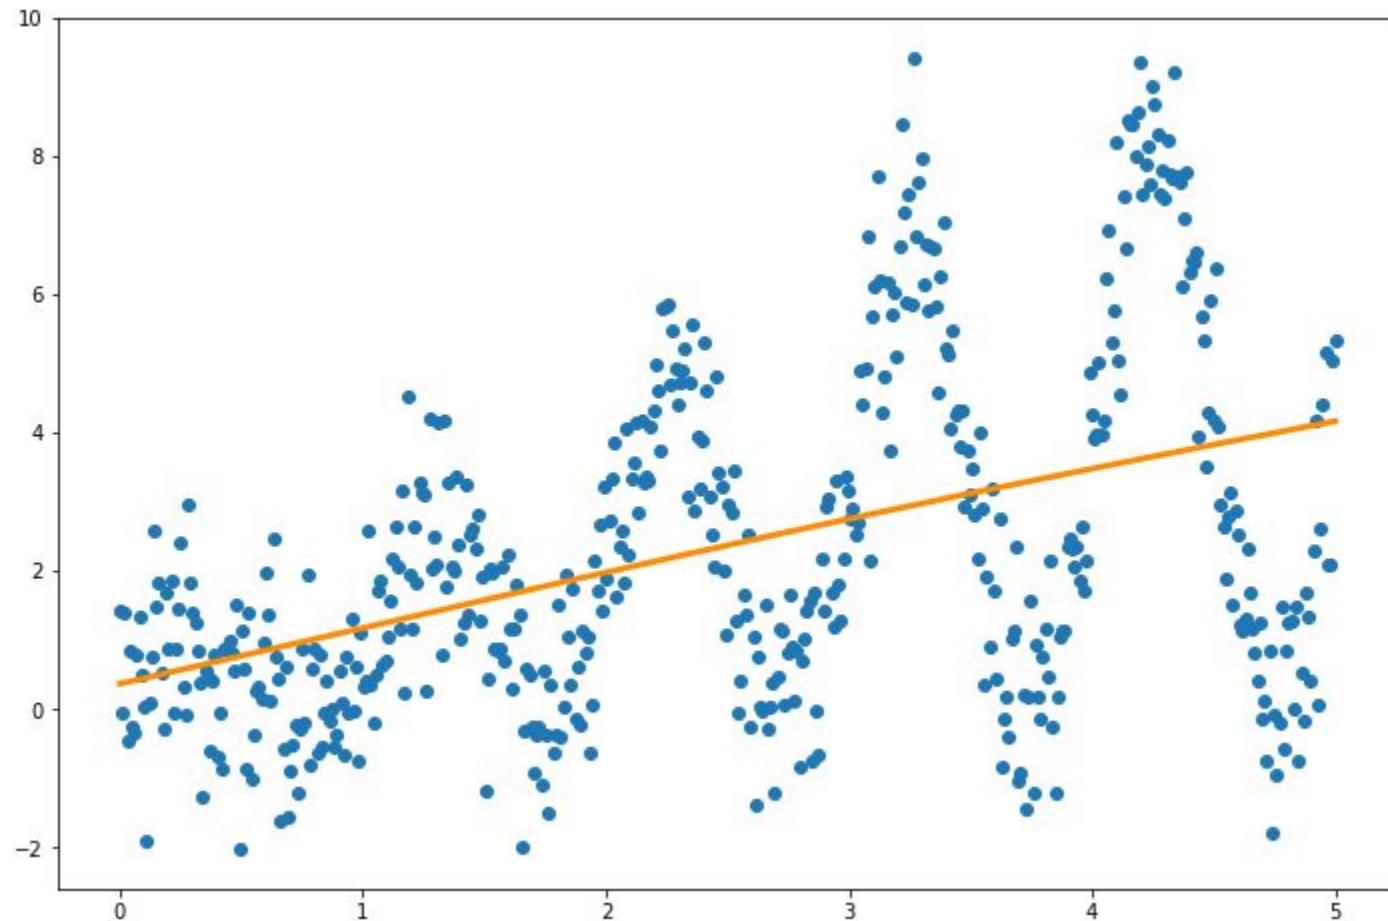
Smoothing

```
from pygam import LinearGAM  
  
# high lambda forces a straight line  
gam = LinearGAM(lam=1e6, n_splines=50).fit(X, y)
```

Smoothing

```
from pygam import LinearGAM  
  
# high lambda forces a straight line  
gam = LinearGAM(lam=1e6, n_splines=50).fit(X, y)
```

Smoothing



```
# generate a nice grid  
XX = gam.generate_X_grid()  
plt.plot(XX, gam.predict(XX))
```

Smoothing

```
>> gam.summary()
```

Smoothing

LinearGAM

Distribution:

NormalDist

Effective DoF: 1.1141

Link Function:

IdentityLink

Log Likelihood: 1339.885

Number of Samples:

500

AIC:

2685.9981

AICc:

2686.0498

GCV:

5.3358

Scale:

5.2952

Pseudo R-Squared:

0.2076

Feature	Function	Data Type	Num Splines	Spline Order	Linear Fit	Lambda	P > x	Sig.	Code
feature 1		numerical	50	3	False	1000000.0	1.11e-16	***	
intercept								1.11e-16	***

Significance codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

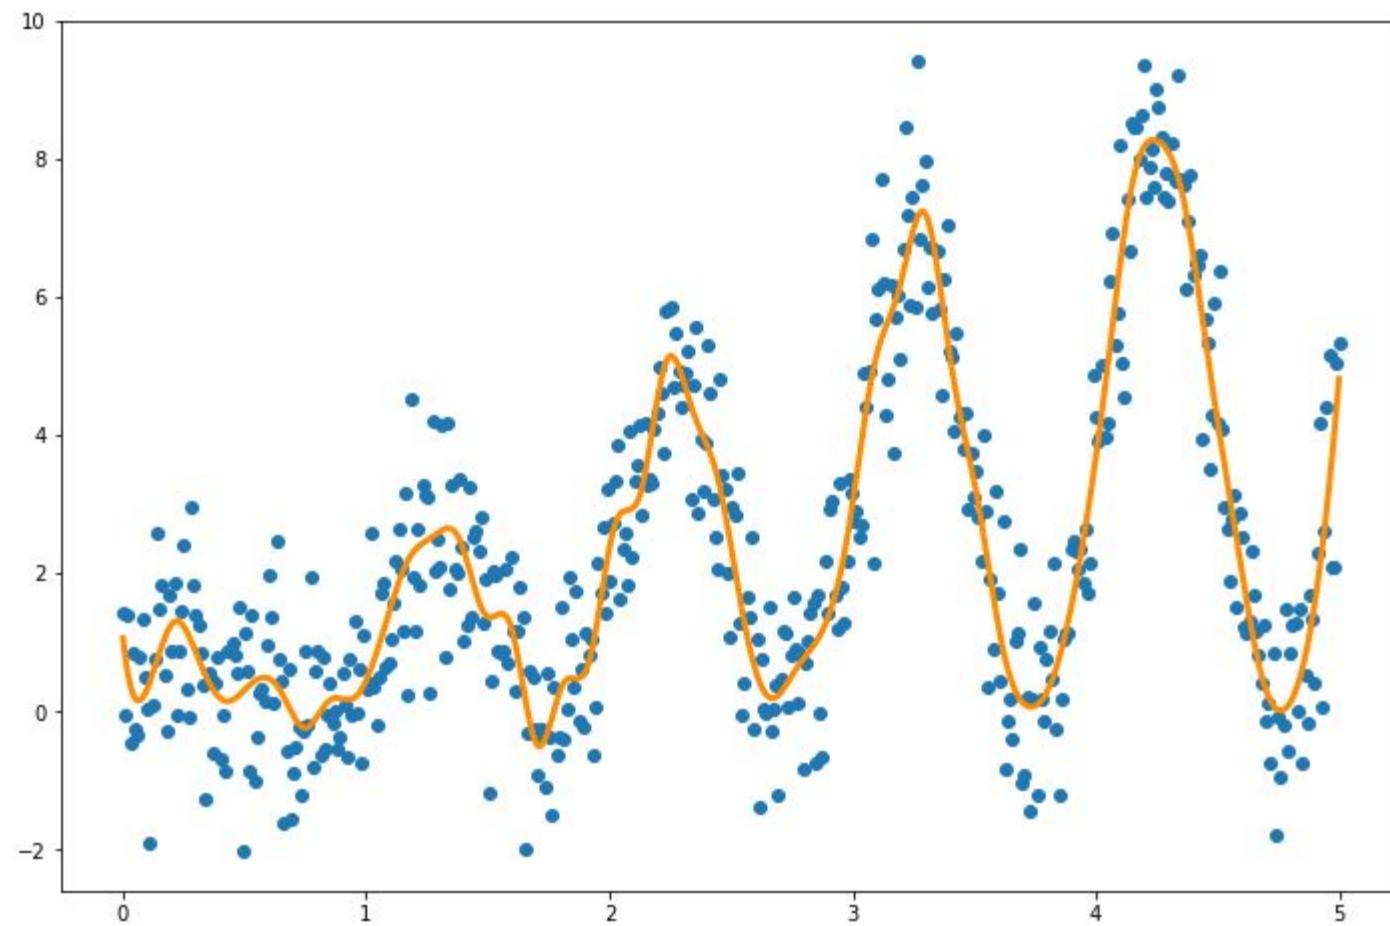
WARNING: Fitting splines and a linear function to a feature introduces a model identifiability problem which can cause p-values to appear significant when they are not.

WARNING: p-values calculated in this manner behave correctly for un-penalized models or models with known smoothing parameters, but when smoothing parameters have been estimated, the p-values are typically lower than they should be, meaning that the tests reject the null too readily.

Smoothing

```
# tiny lambda is super flexible  
gam = LinearGAM(lam=1e-6, n_splines=50).fit(X, y)
```

Smoothing



Smoothing

LinearGAM

Distribution:	NormalDist	Effective DoF:	48.9993
Link Function:	IdentityLink	Log Likelihood:	674.2603
Number of Samples:	500	AIC:	1450.5194
		AICc:	1462.3583
		GCV:	1.1718
		Scale:	0.963
		Pseudo R-Squared:	0.8616

Feature	Function	Data Type	Num Splines	Spline Order	Linear Fit	Lambda	P > x	Sig.	Code
feature 1		numerical	50	3	False	0.0	1.11e-16	***	
intercept							1.11e-16	***	
Significance codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1									

WARNING: Fitting splines and a linear function to a feature introduces a model identifiability problem which can cause p-values to appear significant when they are not.

WARNING: p-values calculated in this manner behave correctly for un-penalized models or models with known smoothing parameters, but when smoothing parameters have been estimated, the p-values are typically lower than they should be, meaning that the tests reject the null too readily.

Smoothing

So how do we choose the ****right**** amount of smoothing?

- Cross validation!

Smoothing

pyGAM computes GCV

- Fast approximation to Leave-One-Out CV

$$\frac{n||y - \hat{\mu}||}{(n - tr(A))^2}$$

Smoothing

pyGAM computes GCV

- Fast approximation to Leave-One-Out CV

$$\frac{n \left| \left| \mathbf{y} - \hat{\boldsymbol{\mu}} \right| \right|}{\left(n - edof \right)^2}$$

Smoothing

```
gam = LinearGAM(n_splines=50).fit(X, y)
```

Smoothing

```
gam = LinearGAM(n_splines=50).gridsearch(X, y)
```

Smoothing

```
lams = np.logspace(-4, 0, 100)
gam = LinearGAM(n_splines=50).gridsearch(X, y, lam=lams)
```

Smoothing

LinearGAM

Distribution:

NormalDist Effective DoF:

25.1906

Link Function:

IdentityLink Log Likelihood:

692.0659

Number of Samples:

500 AIC:

1438.5129

AICc:

1441.7621

GCV:

1.0855

Scale:

0.9836

Pseudo R-Squared:

0.8477

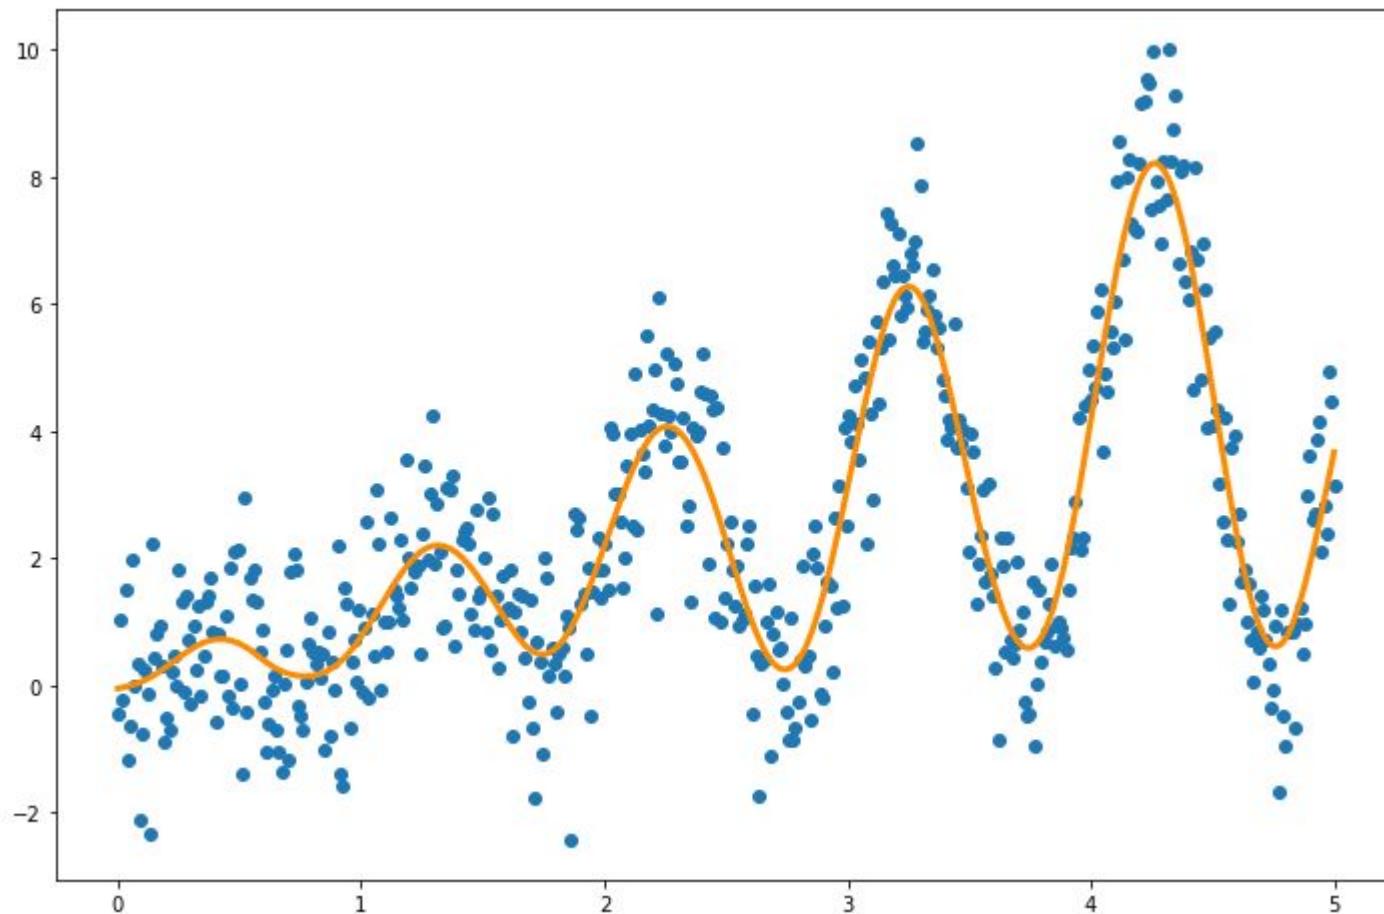
Feature	Function	Data Type	Num Splines	Spline Order	Linear Fit	Lambda	P > x	Sig.	Code
feature 1		numerical	50	3	False	1.0	1.11e-16	***	
intercept							1.11e-16	***	

Significance codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

WARNING: Fitting splines and a linear function to a feature introduces a model identifiability problem which can cause p-values to appear significant when they are not.

WARNING: p-values calculated in this manner behave correctly for un-penalized models or models with known smoothing parameters, but when smoothing parameters have been estimated, the p-values are typically lower than they should be, meaning that the tests reject the null too readily.

Smoothing



Back at PyDatatopolis

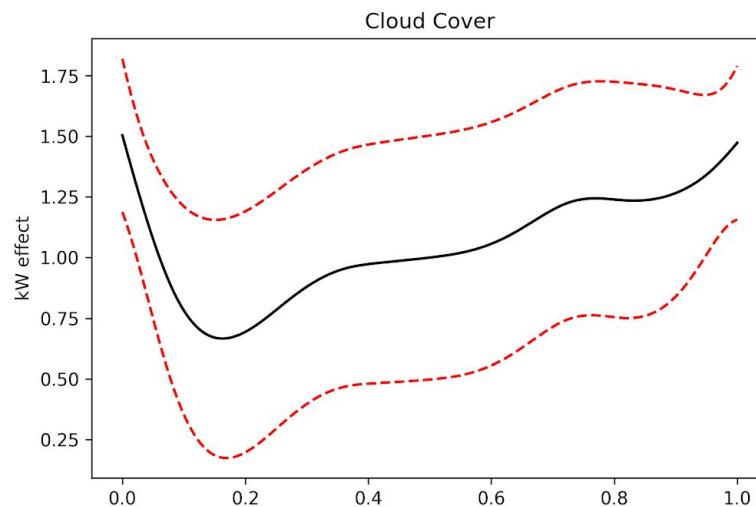
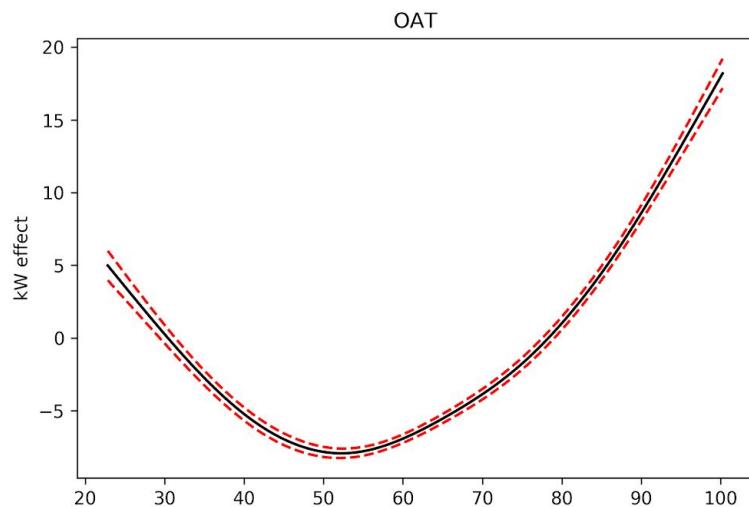
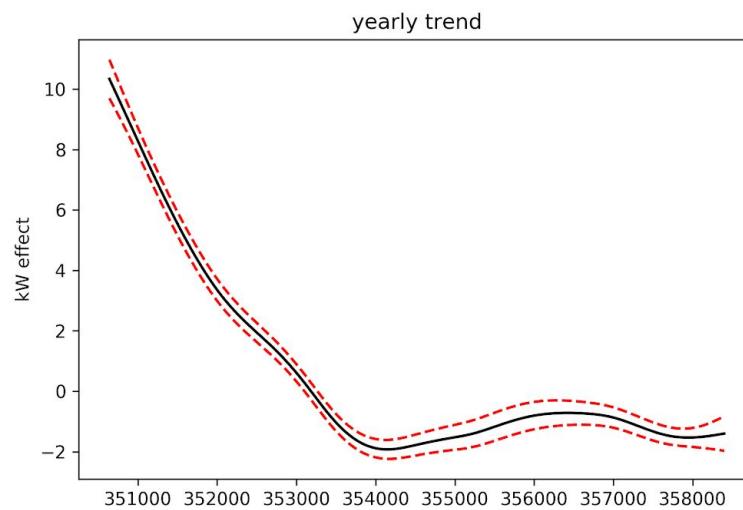
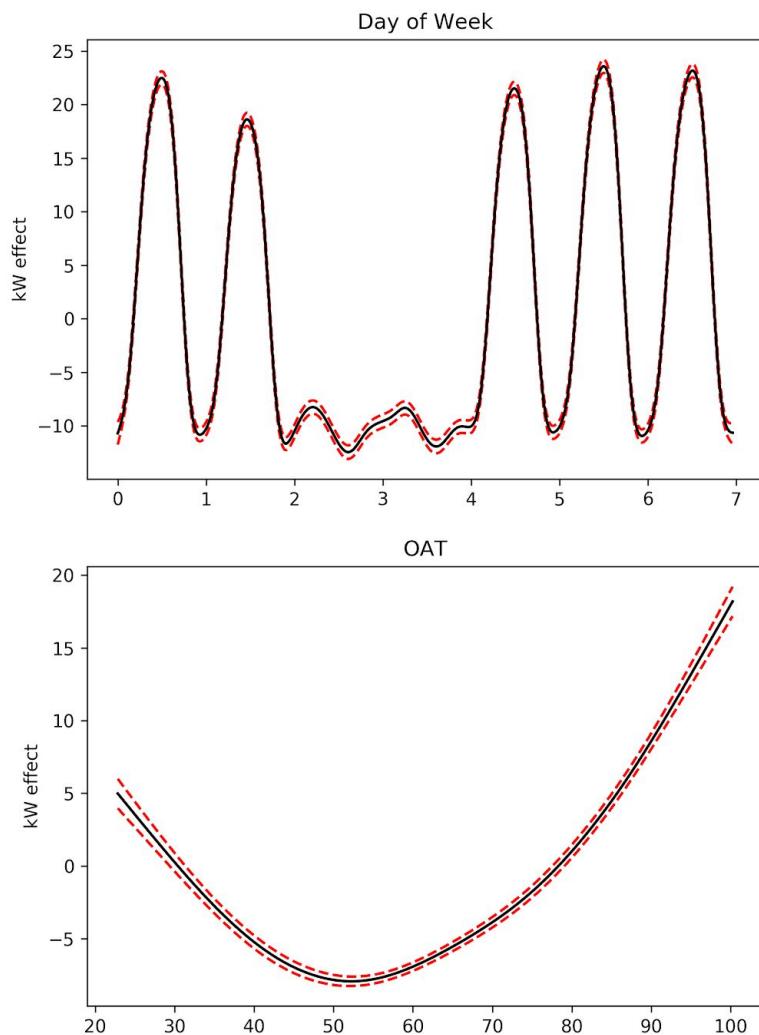
- Select right amount of smoothing
- Cloud cover?

Back at PyDatatopolis

```
dist = np.random.uniform(1e-3, 1e3, (100,4))

gam = LinearGAM(n_splines=[100,25,25,25]).gridsearch(X, y, lam=dist)
```

Back at PyDatatopolis



Back at PyDatatopolis

LinearGAM

Distribution:	NormalDist	Effective DoF:	56.6547					
Link Function:	IdentityLink	Log Likelihood:	-32515.6534					
Number of Samples:	7759	AIC:	65146.6162					
		AICc:	65147.4945					
		GCV:	26.2008					
		Scale:	25.8566					
		Pseudo R-Squared:	0.8865					
Feature	Function	Data Type	Num Splines	Spline Order	Linear Fit	Lambda	P > x	Sig. Code
feature 1		numerical	100	3	False	36.4721	1.11e-16	***
feature 2		numerical	25	3	False	229.9264	1.11e-16	***
feature 3		numerical	25	3	False	938.4219	1.11e-16	***
feature 4		numerical	25	3	False	430.7849	4.70e-01	
intercept							1.11e-16	***
Significance codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1								

Future work

Lots of cool features...

- Feature interactions
- Periodicity constraints
- Direct smoothing parameter optimization
- ExpectileGAM
- ...

<https://github.com/dswah/pyGAM/issues>

Vielen Dank!



Questions?

- Can pyGAM do classification?
- What other distributions can we model?
- What size of data can pyGAM handle?
- Patsy compatible?

Some math...

